



Monitoring with Prometheus

Key Takeaways

Introduction to Prometheus

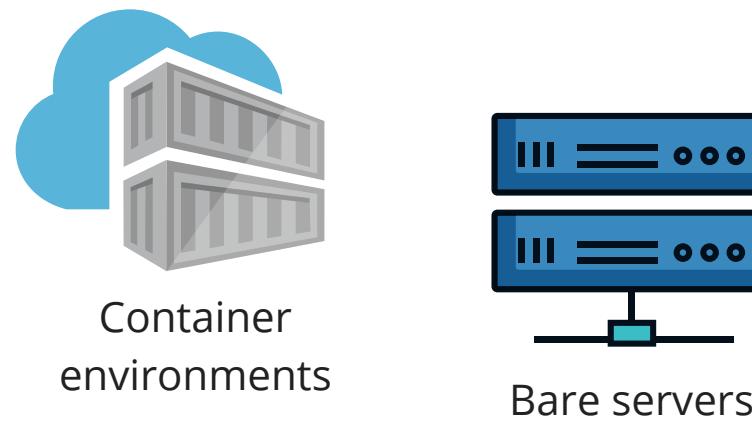
- Prometheus is an **open-source monitoring** system and **alerting** toolkit
- Prometheus is used widely and has an active community
- It gathers, organizes, and **stores metrics as time series data from targets** by "scraping" metrics HTTP endpoints
- Can trigger alerts when specified conditions are observed



Why we need a monitoring tool - 1

Visibility in different environments

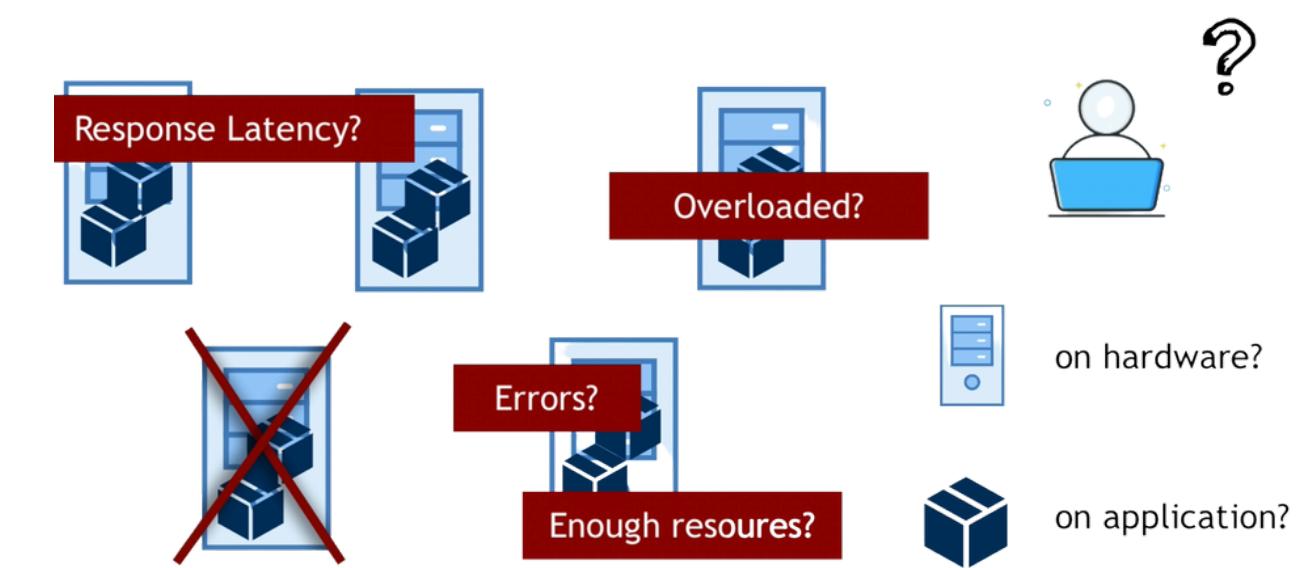
- You need **visibility in all kinds of environments**, but especially in a highly dynamic container environment, which is more challenging to monitor



- For that - you need tools like Prometheus, which are designed for monitoring these types of environments

Visibility on different levels

- When you have 100s or 1000s of containers, plus components on multiple levels (infrastructure, platform, application) you need a way to have a visibility and **consistent monitoring across all these components**



- Without visibility, it's a black box for you. When things break inside your complex environment, you have **no idea what is happening**. You don't know what has caused the issue, what is not working.

Why we need a monitoring tool - 2

Example Use Case

Problem:

- Many application errors appear on the front-end to the end user
- They **only see the error message**, but the cause can be any of the many components in the back-end

Solution:

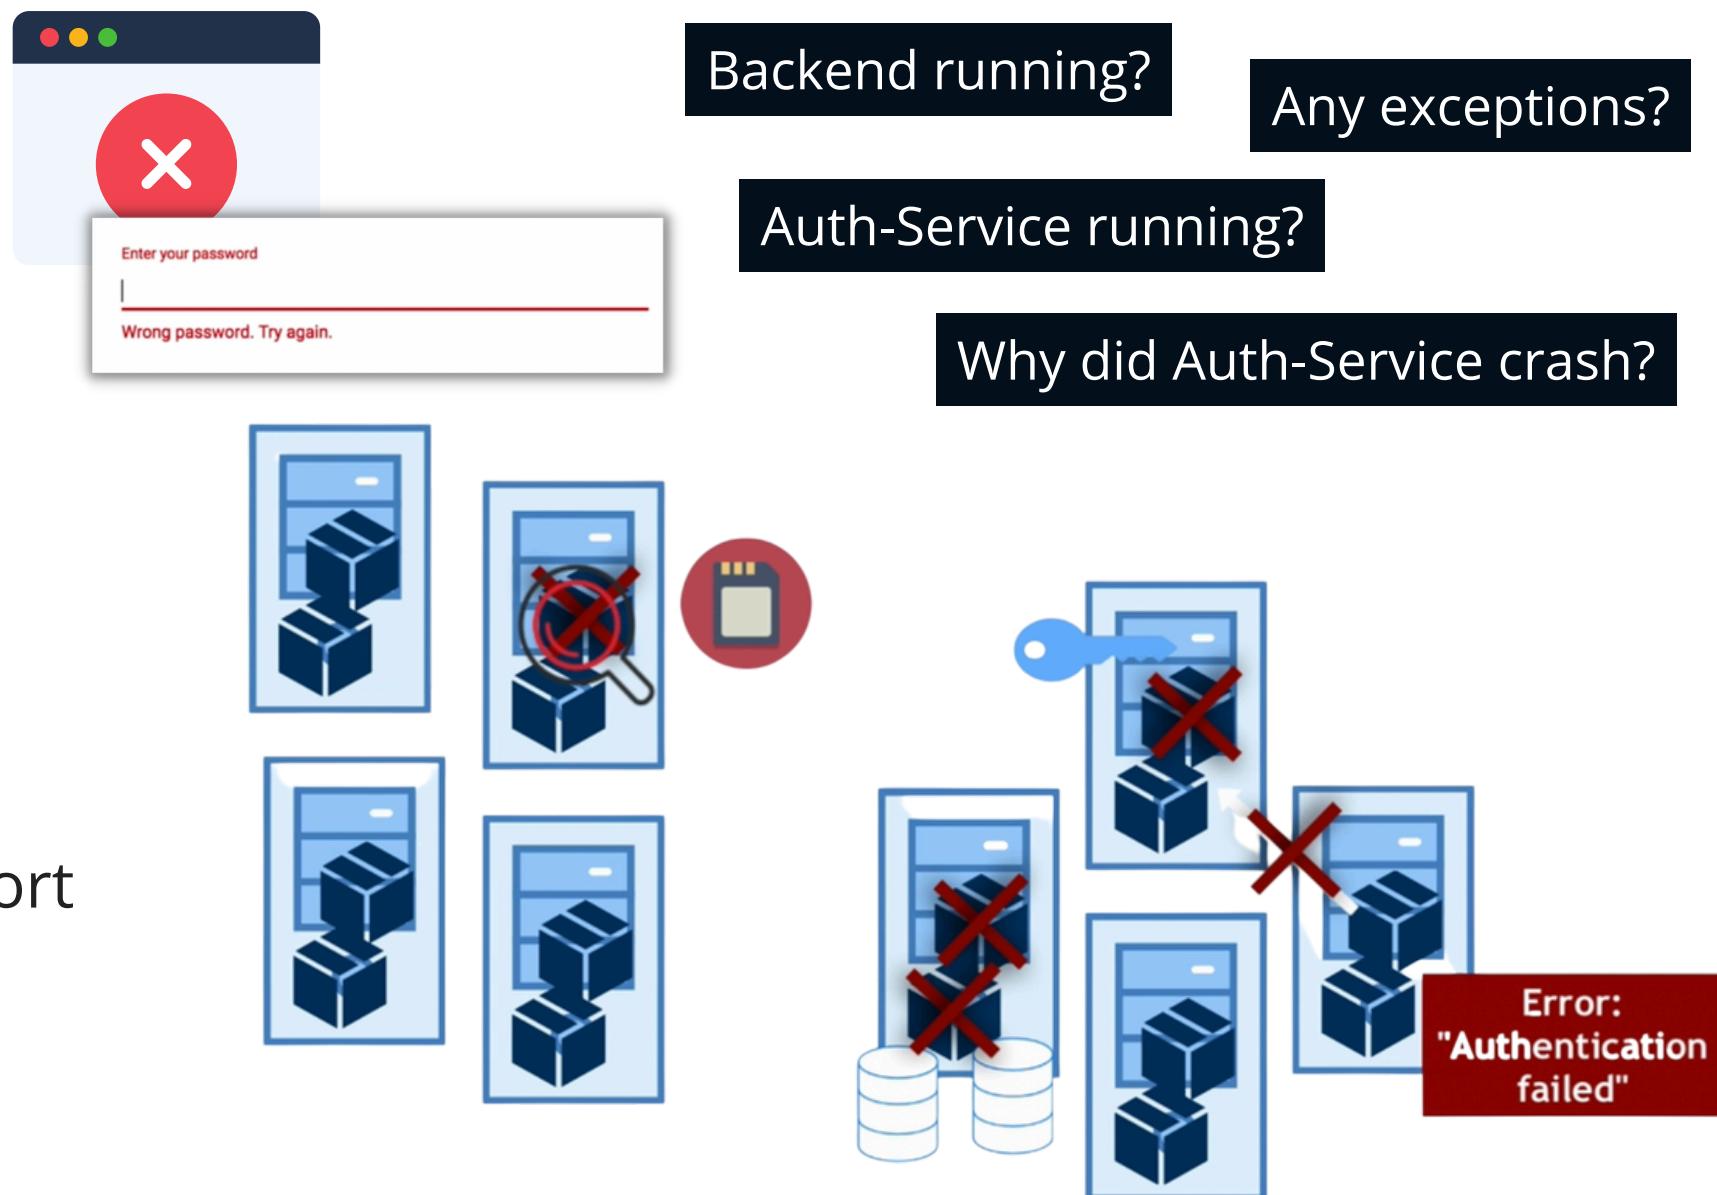
- Monitoring can help identify the problem quickly with little effort
- Instead of manually trying to troubleshoot across multiple components, it will **help pin point directly to the root cause**



Saves you a lot of time and effort



Company saves a lot of money



Why we need a monitoring tool - 3

With Prometheus, everything is automated. It's constantly monitoring and looking out for any issues in real time and may even identify a potential issue before it happens, so you can prevent it.



Constantly monitors all the services



Triggers alerts when a service crashes



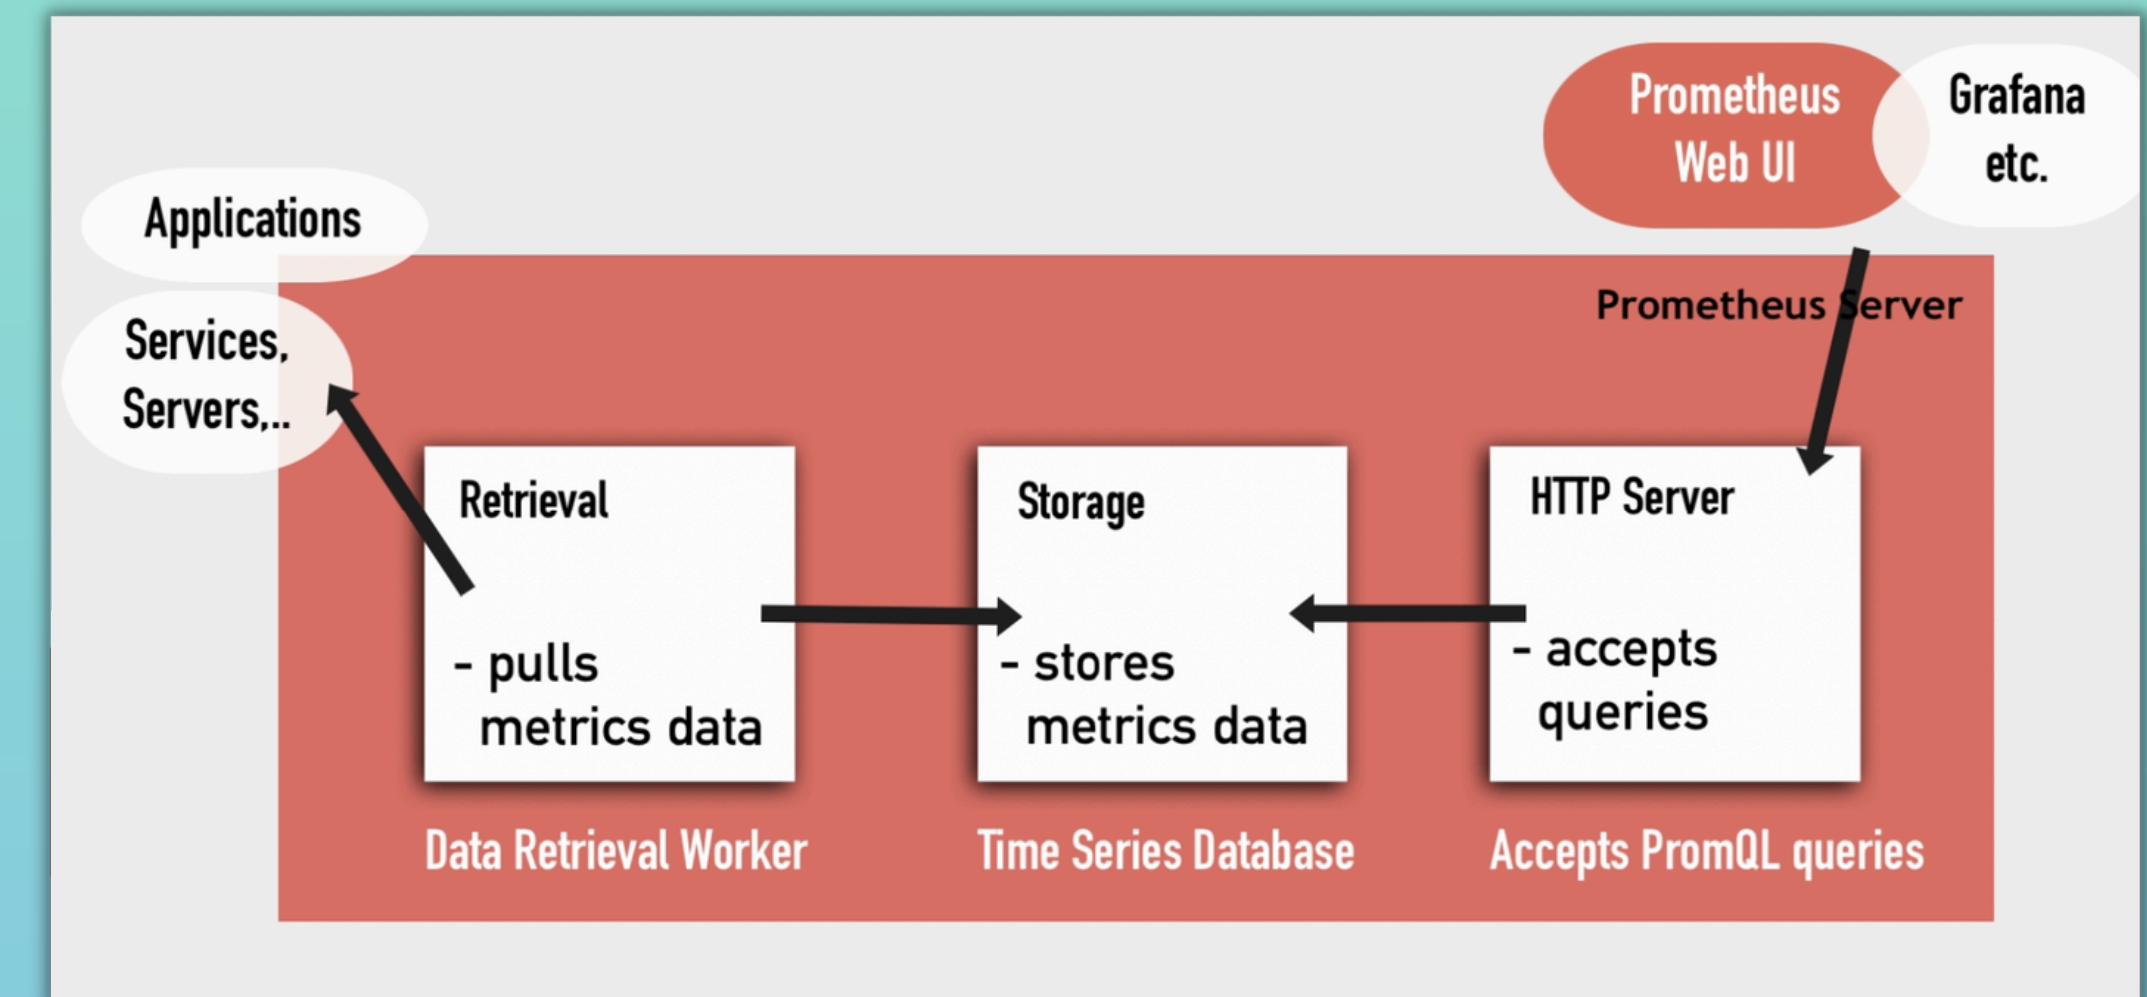
Helps to identify problems before they happen

Prometheus Architecture - 1

How it all works

Prometheus Server

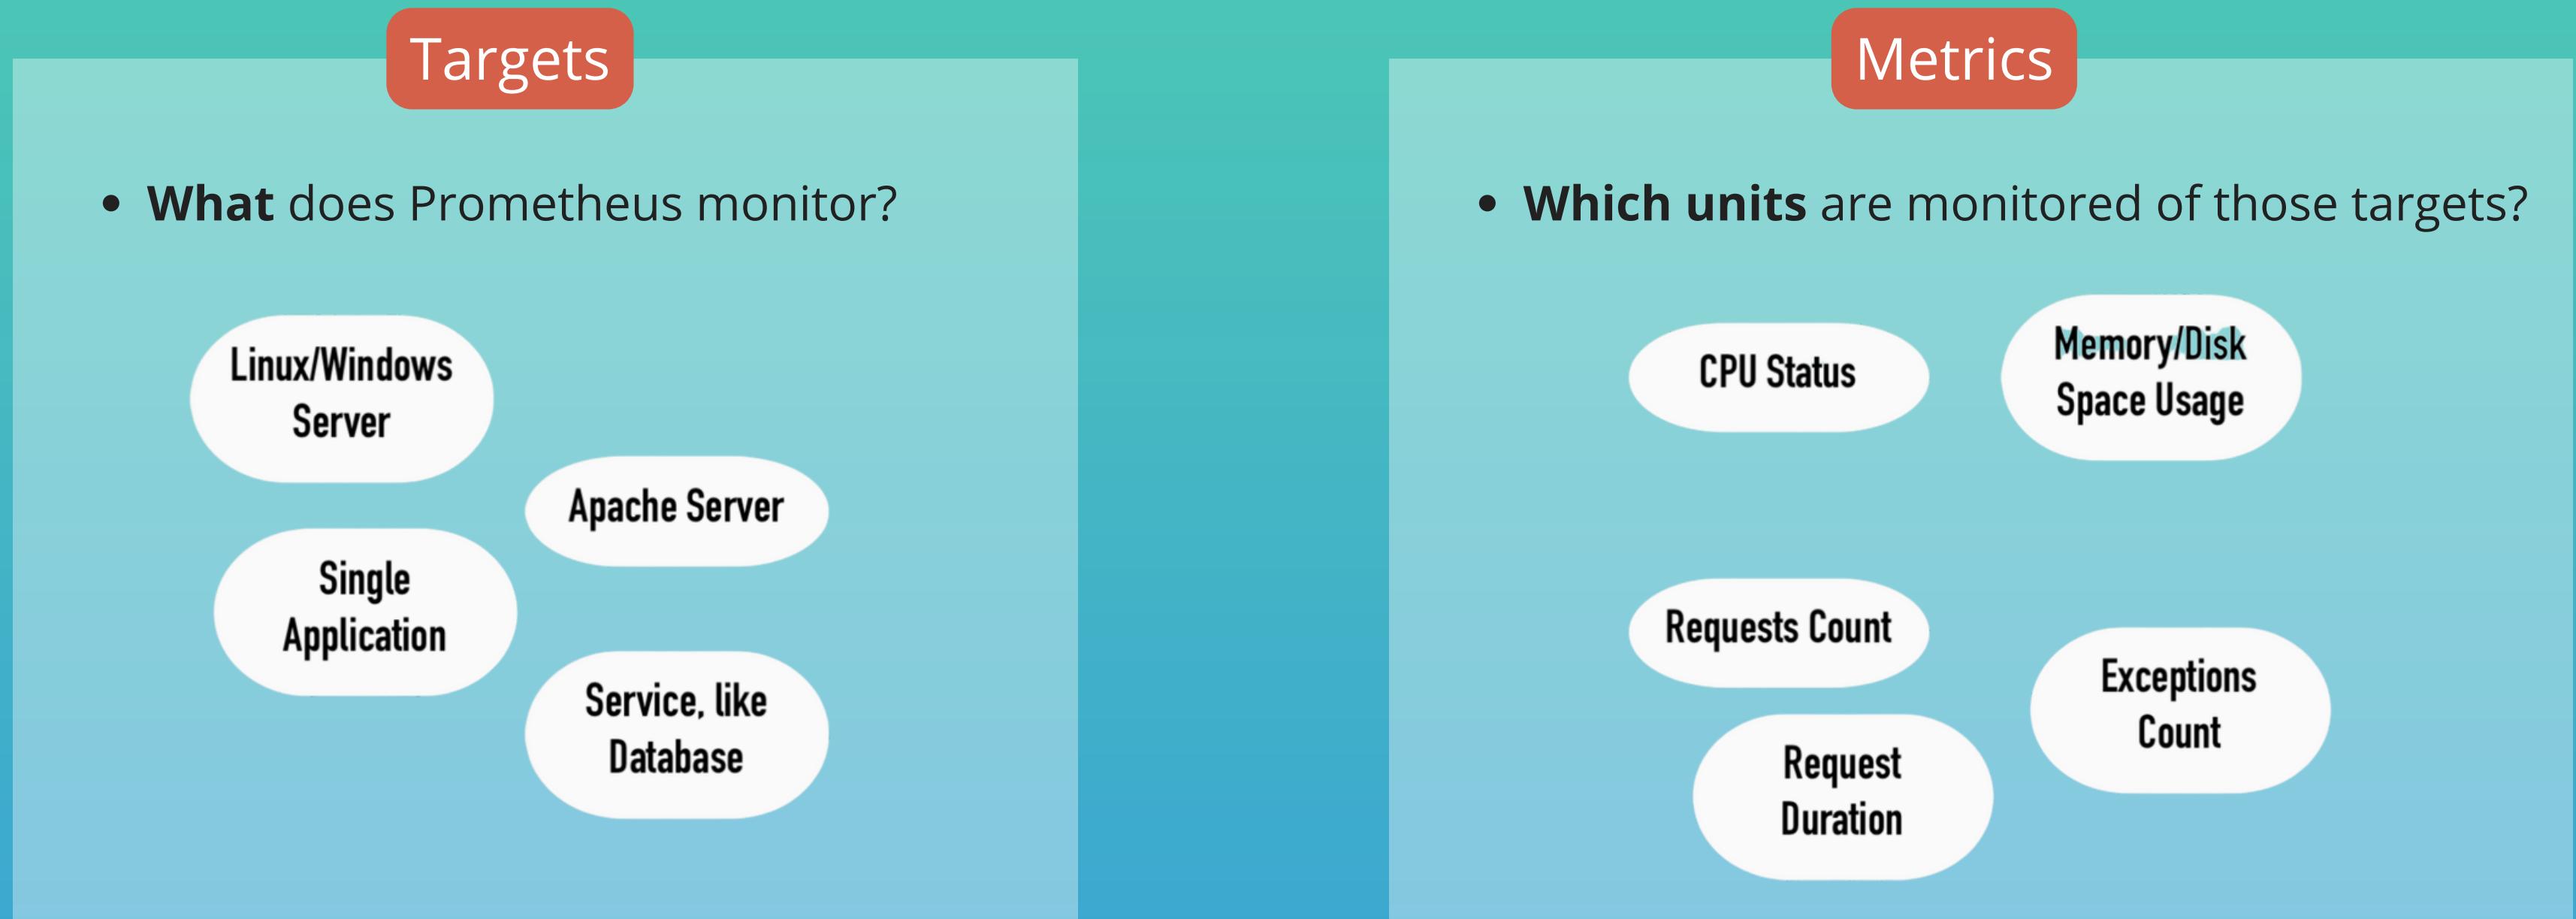
- Is the **main component**
- Does the actual monitoring work
- Scrapes and stores time series data



Prometheus Architecture - 2

How it all works: Targets & Metrics

- Prometheus pulls metrics from targets



Prometheus Architecture - 3

How it all works: Metrics

- Metrics play an important role in understanding why your application is working in a certain way

Metric Entries

- Format: Human-readable text-based
- Metric entries consist of: **TYPE** and **HELP** attributes

HELP ↗ description of what the metrics is

TYPE ↗ 3 metrics types

1) Counter

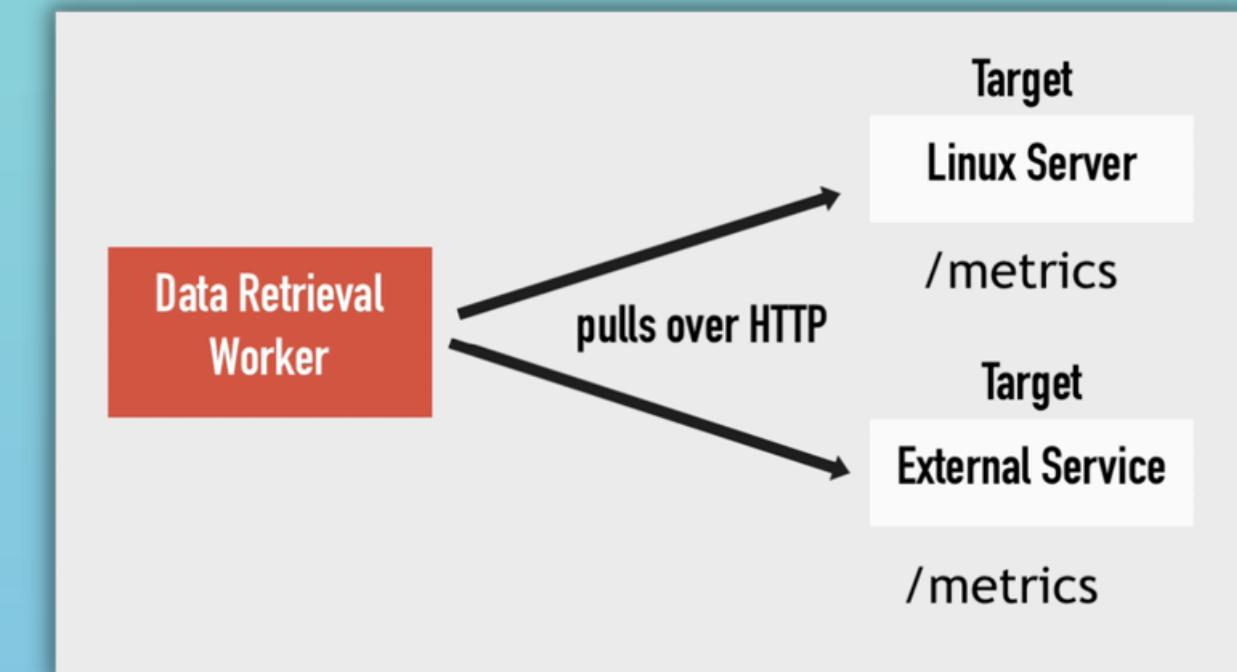
2) Gauge

3) Histogram

...how many times x happened ...what is current value of x now? ...how long or how big?

How Prometheus collects Metrics Data from Targets

- Prometheus **pulls from HTTP endpoints**
- Targets must expose: [hostaddress]/**metrics**
- Must be in correct format that Prometheus understands



Prometheus Architecture - 4

How it all works: Exporters

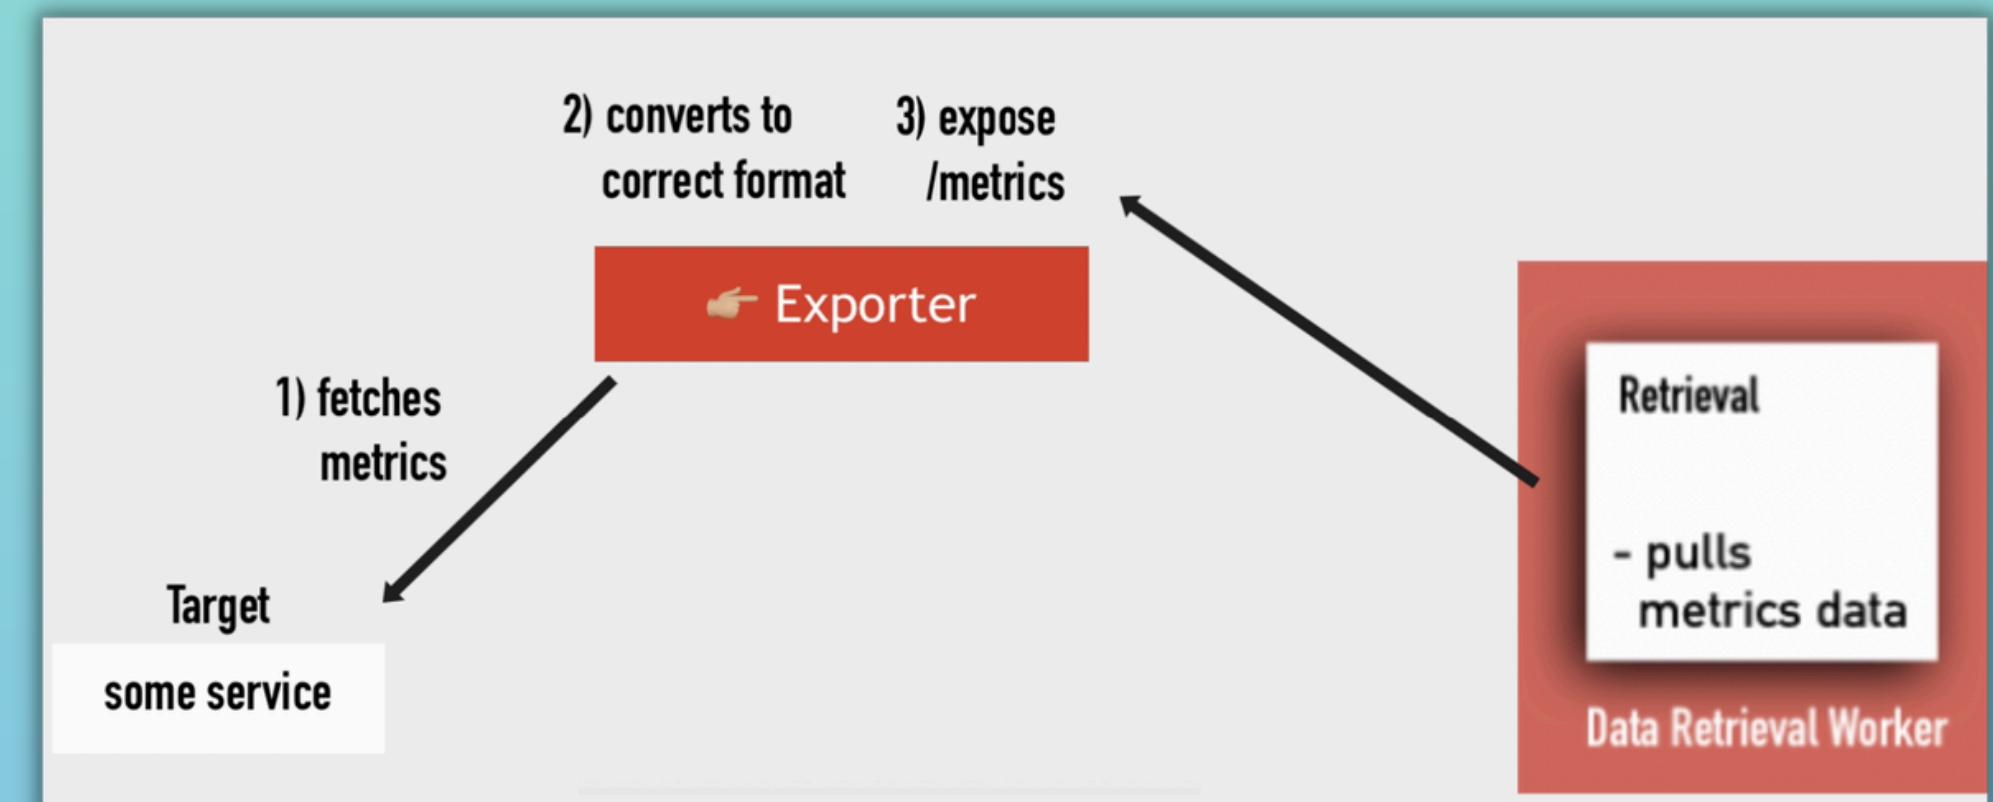
- Some services expose `/metrics` endpoints by default
- Others need another component for that:

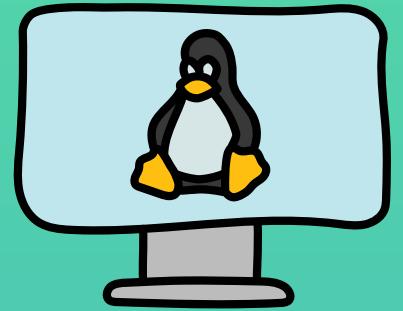
Exporters

- Exporters help in exporting existing metrics from third-party systems as Prometheus metrics
- An exporter is a service that **fetches metrics from a target and converts the data and exposes them as Prometheus metrics**
- Prometheus can then scrape this endpoint as usual

Official vs Third-Party

- Some are maintained as part of the official Prometheus organization, others are externally contributed and maintained.





Prometheus Architecture - 5

How it all works: Exporters

Example: Monitor a Linux Server

1. Download a **node exporter**
2. Untar and execute
3. Converts metrics of the server
4. Exposes /metrics endpoint
5. Configure Prometheus to scrape this endpoint

Hardware related

- [apcupsd exporter](#)
- [BIG-IP exporter](#)
- [Collins exporter](#)
- [Dell Hardware OMSA exporter](#)
- [IBM Z HMC exporter](#)
- [IoT Edison exporter](#)
- [IPMI exporter](#)
- [knxd exporter](#)
- [Modbus exporter](#)
- [Netgear Cable Modem Exporter](#)
- [Netgear Router exporter](#)
- [Node/system metrics exporter \(official\)](#)
- [NVIDIA GPU exporter](#)
- [ProSAFE exporter](#)
- [Ubiquiti UniFi exporter](#)



Exporters are available as Docker Images

Example: Monitor own applications

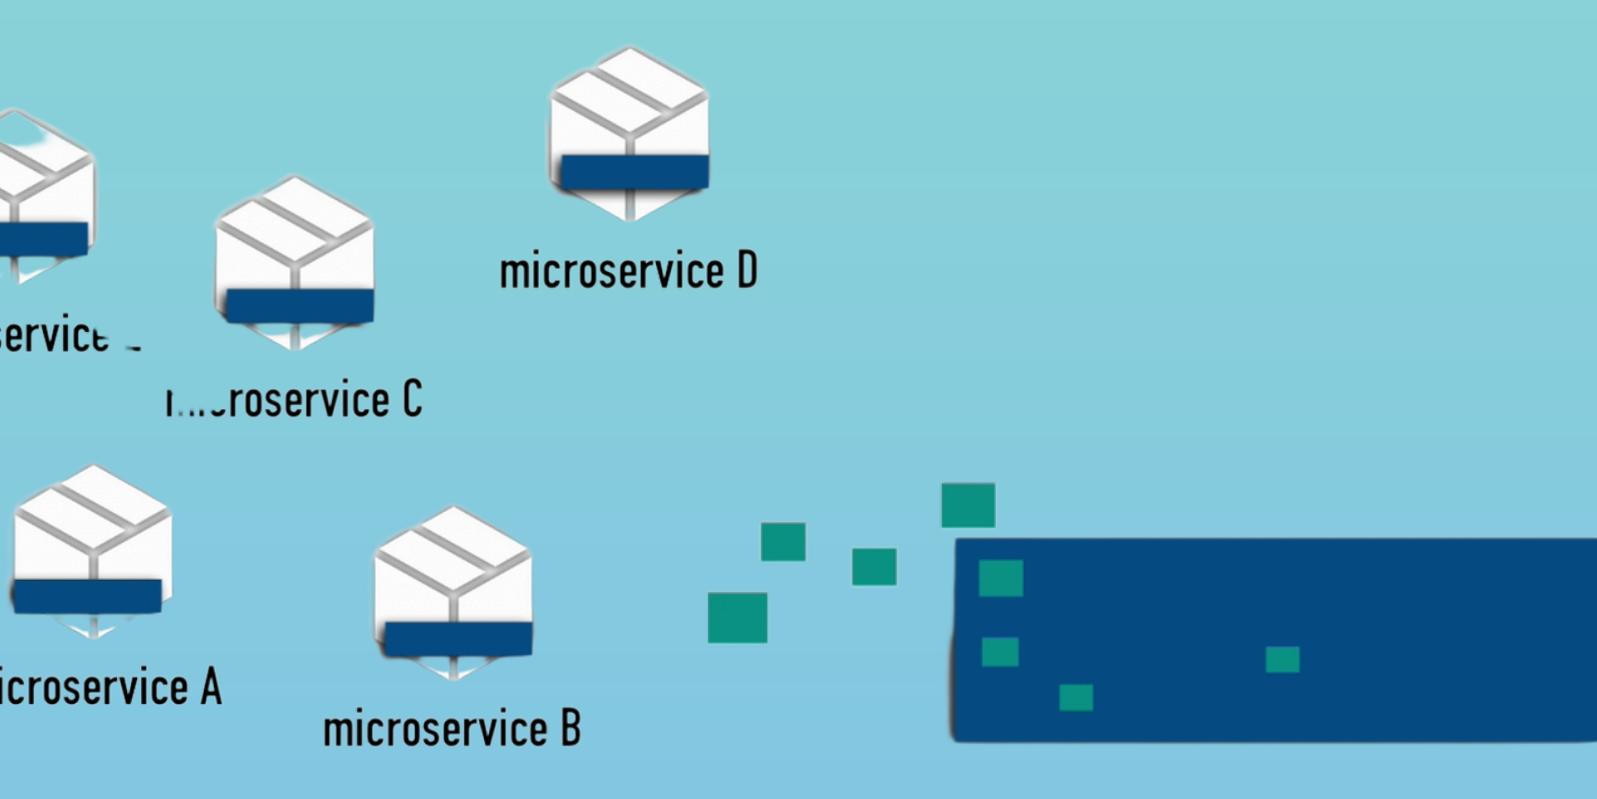
- Client libraries let you **define and expose internal metrics** via an HTTP endpoint on your application's instance
- Metrics like:
 - How many requests?
 - How many exceptions?
 - Server resources used?
- Choose a Prometheus client library that matches the language in which your application is written

Prometheus Architecture - 6

How it all works: Push vs Pull

- **Important difference** of Prometheus compared to other monitoring systems like Amazon Cloud Watch or New Relic

Others - Push Model



- Services **push to a centralized collection platform**

Prometheus - Pull Model

- Prometheus pulls metrics from endpoints

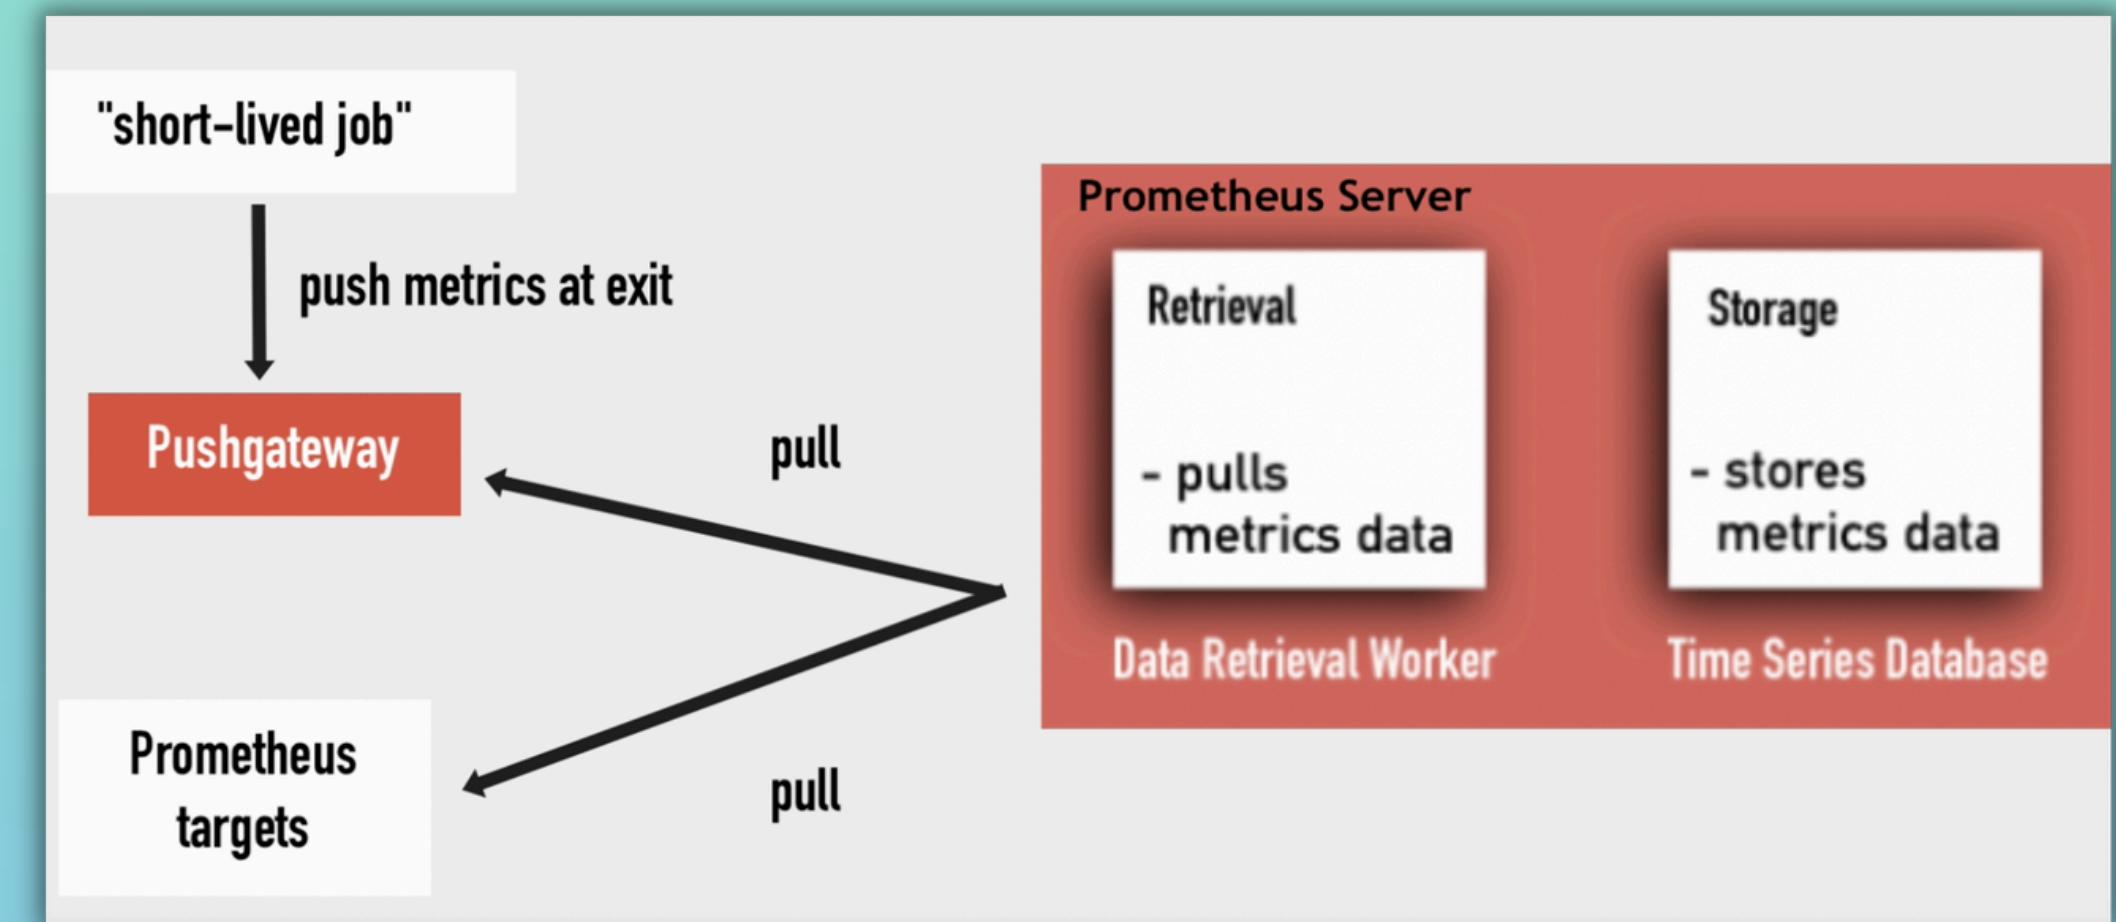
- ✖ High load of network traffic
- ✖ Monitoring can become your bottleneck
- ✖ Installation of additional software to push metrics

Prometheus Architecture - 7

How it all works: Pushgateway

Pushgateway

- An intermediary service, which **allows you to push metrics** from jobs, which cannot be scraped
- Prometheus recommends using the Pushgateway only in certain limited cases: Usually only valid use case for capturing the outcome of a **service-level batch job**



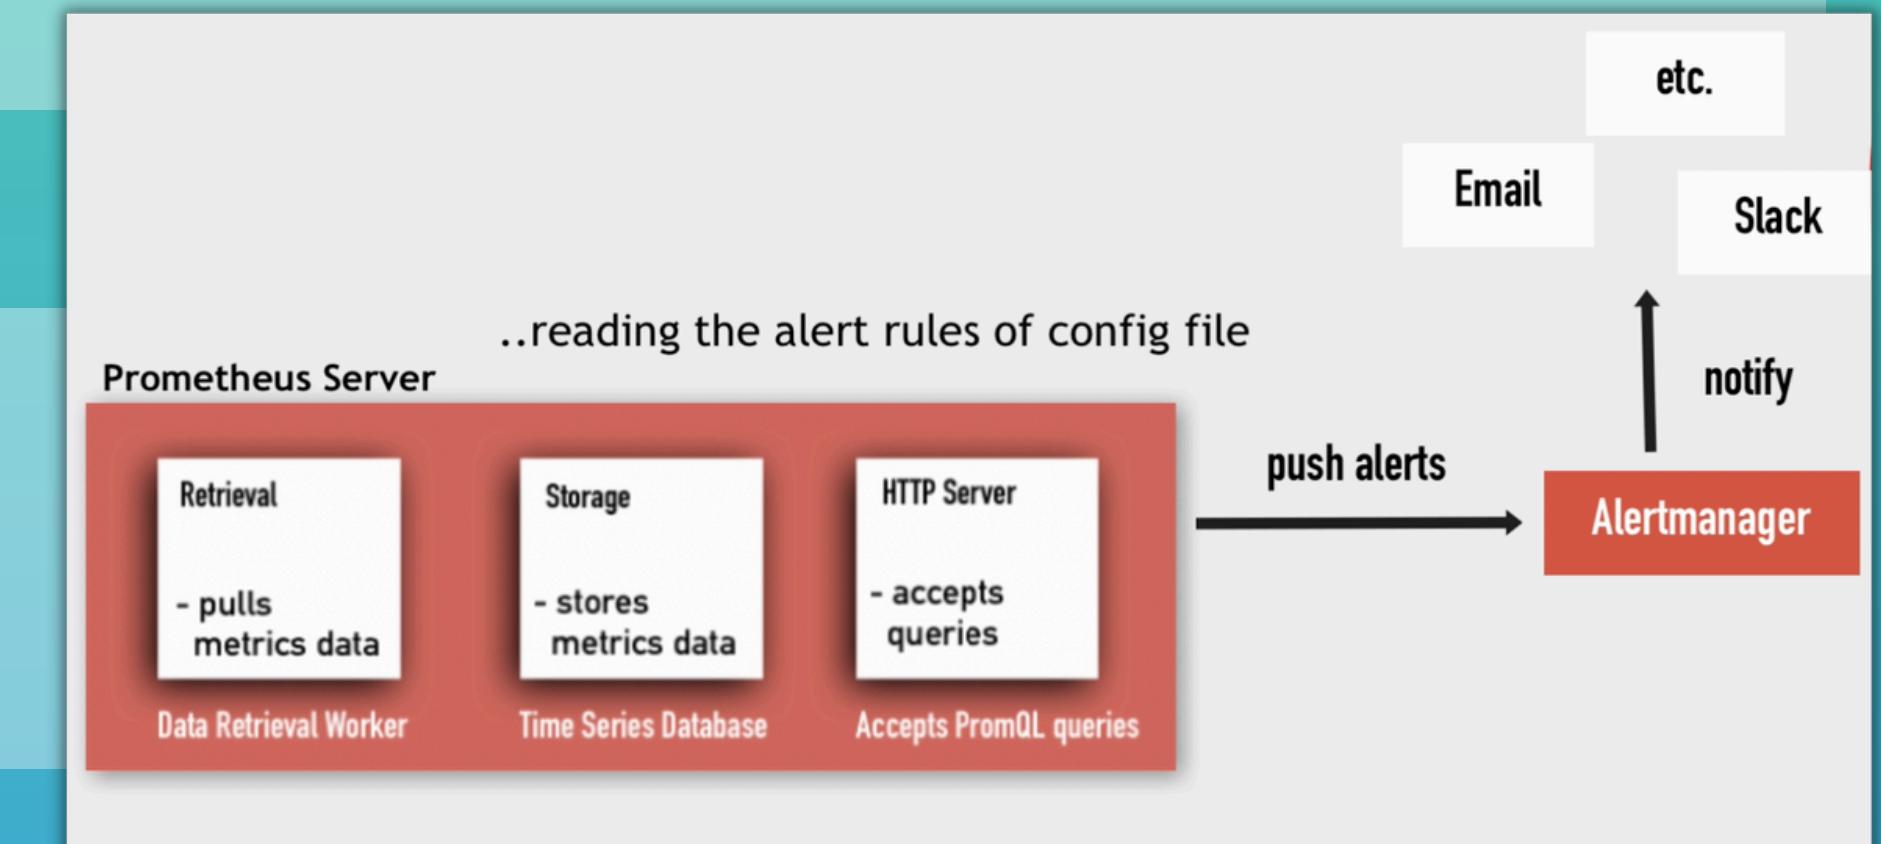
Prometheus Architecture - 8

How it all works: Alertmanager

Alertmanager

- The Alertmanager handles alerts sent by Prometheus server
- Takes care of deduplicating, grouping and routing them to the correct receiver integrations

- Receiver of these alerts can be email, PagerDuty, Slack etc.

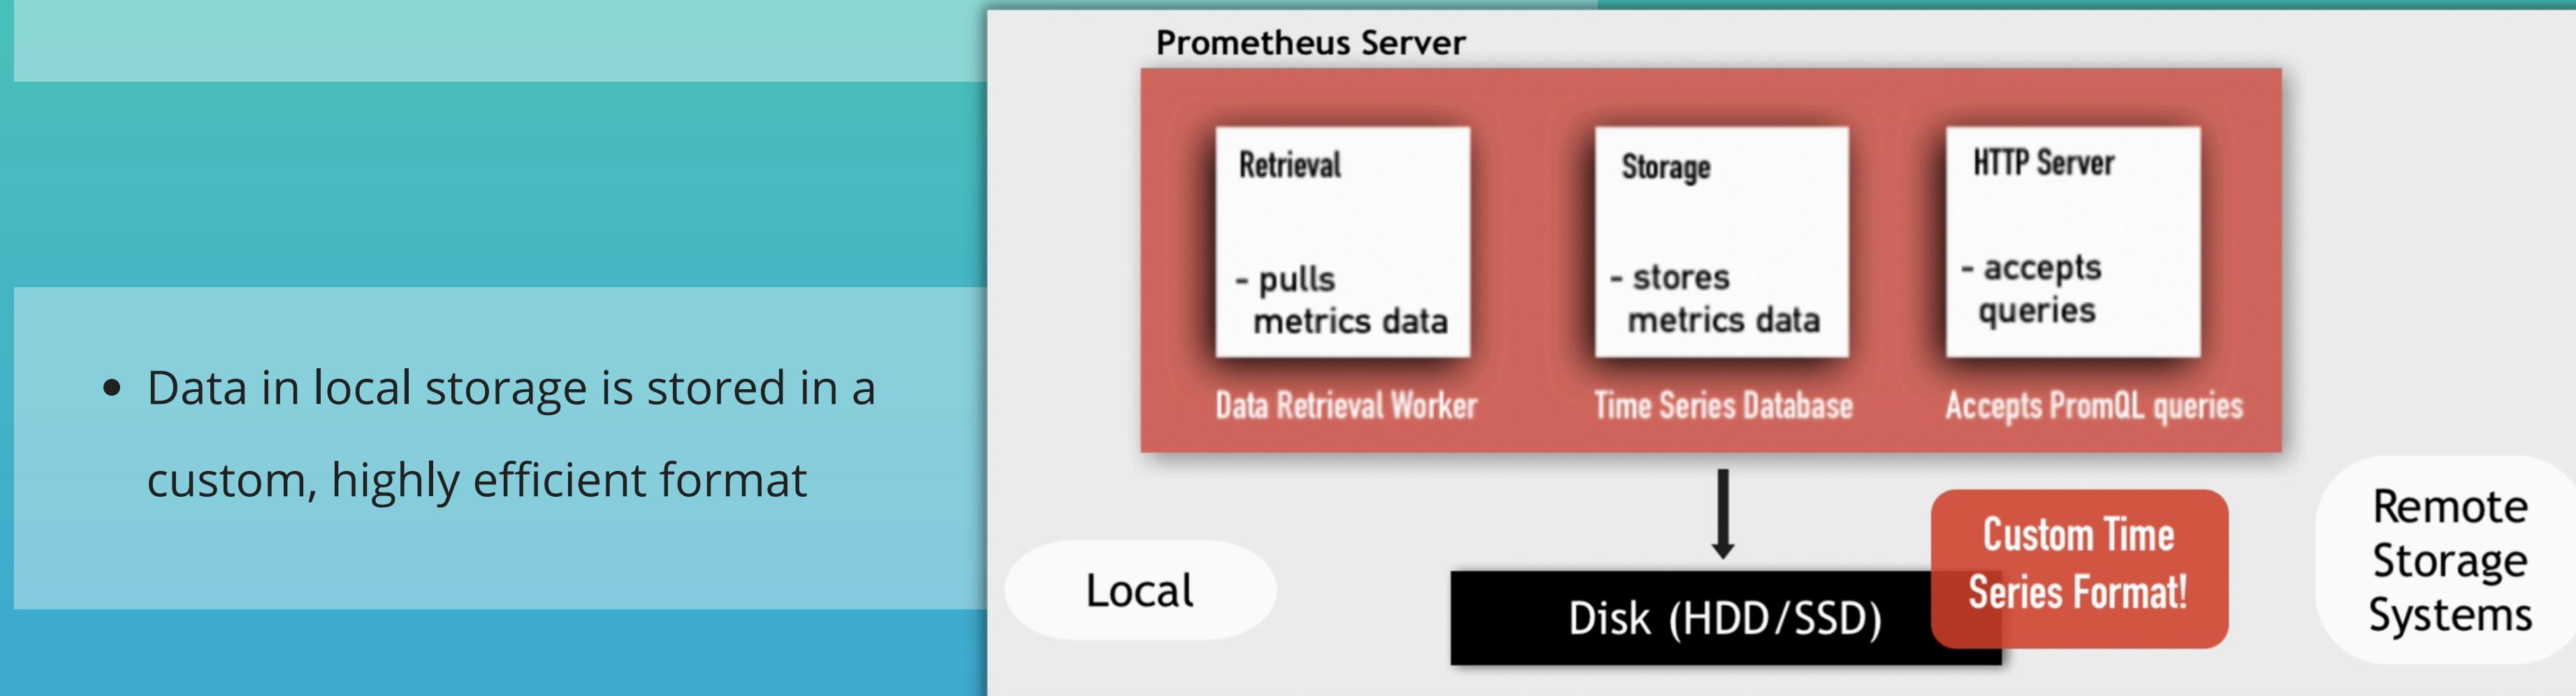
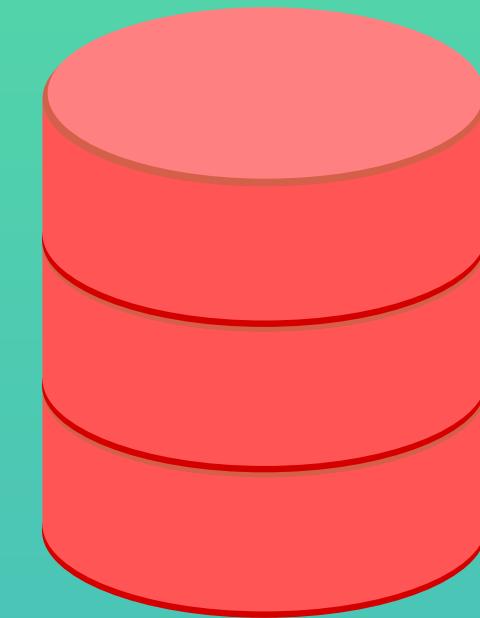


Prometheus Architecture - 9

How it all works: Data Storage

Prometheus Data Storage

- Prometheus includes a local on-disk time series database
- But optionally integrates with remote storage systems

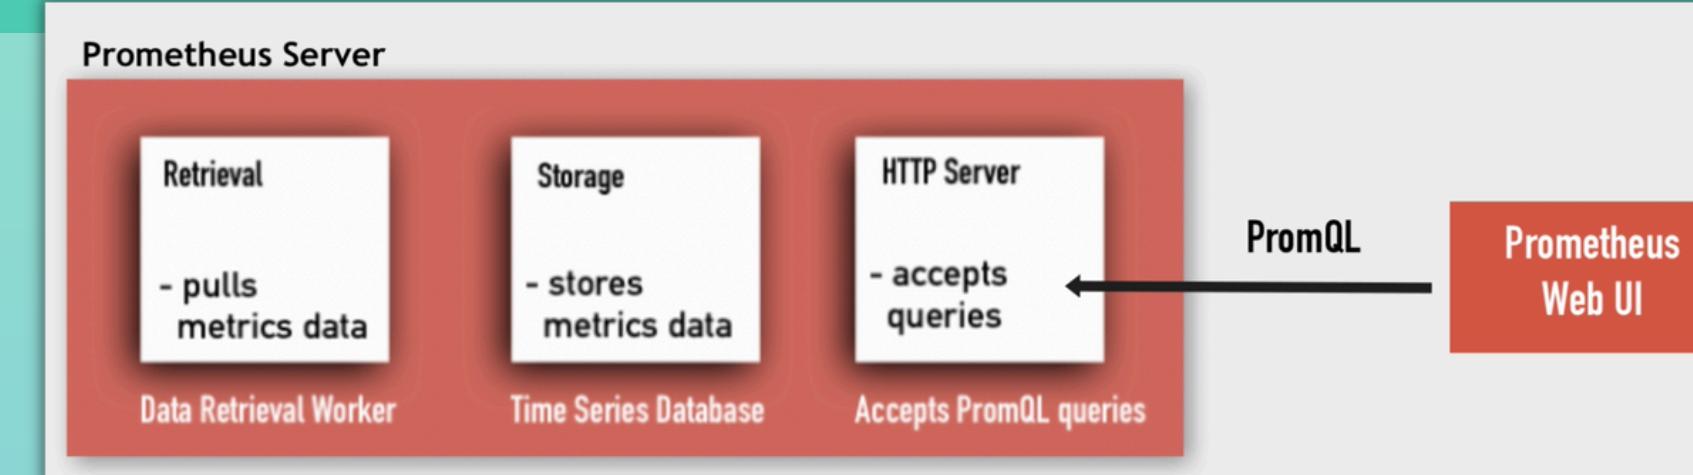


Prometheus Architecture - 10

How it all works: PromQL

Querying Prometheus

- Prometheus provides a functional query language called PromQL
- Let's user select and aggregate time series data in real time



Options to view result

1) Query target directly

2) Prometheus Web UI

3) Or use a more powerful visualization tool, e.g. Grafana



Example Queries:

```
http_requests_total{status!~"4.."}  
↳ Query all HTTP status codes except 4xx ones
```

```
rate(http_requests_total[5m])[30m:]  
↳ Returns the 5-minute rate of the http_requests_total metric for the past 30mins
```

Configuring Prometheus - 1

YAML Config

- You write your configuration in a prometheus.yaml file
- Let Prometheus know **what to scrape and when**:

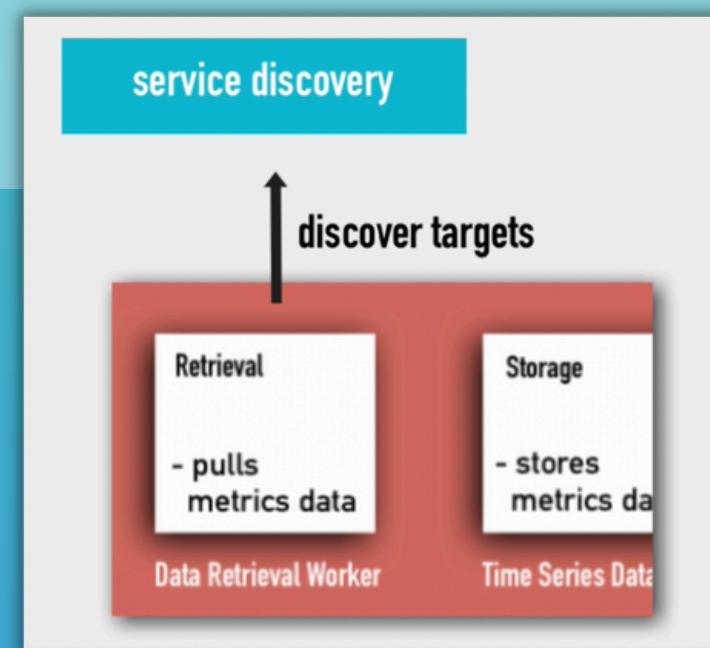
Which targets?

At what interval?



prometheus.yaml

- Targets are discovered via a service discovery mechanism



Example Config File

- **How often** Prometheus will scrape its targets
- **Rules** for aggregating metric values or creating alerts when condition met
- **What resources** Prometheus monitors

```
global:  
  scrape_interval: 15s  
  evaluation_interval: 15s  
  
rule_files:  
  # - "first.rules"  
  # - "second.rules"  
  
scrape_configs:  
  - job_name: prometheus  
    static_configs:  
      - targets: ['localhost:9090']
```

Prometheus has its own
/metrics endpoint

Configuring Prometheus - 2

Define your own jobs

- Default values for each job:

```
metrics_path: "/metrics"
scheme: "http"
```

```
global:
  scrape_interval:      15s
  evaluation_interval: 15s

rule_files:
  # - "first.rules"
  # - "second.rules"

scrape_configs:
  - job_name: "prometheus"
    static_configs:
      - targets: ['localhost:9090']

  - job_name: node_exporter
    scrape_interval: 1m
    scrape_timeout: 1m
    static_configs:
      - targets: ['localhost:9100']
```

Prometheus Characteristics

 Reliable

 Standalone and self-containing

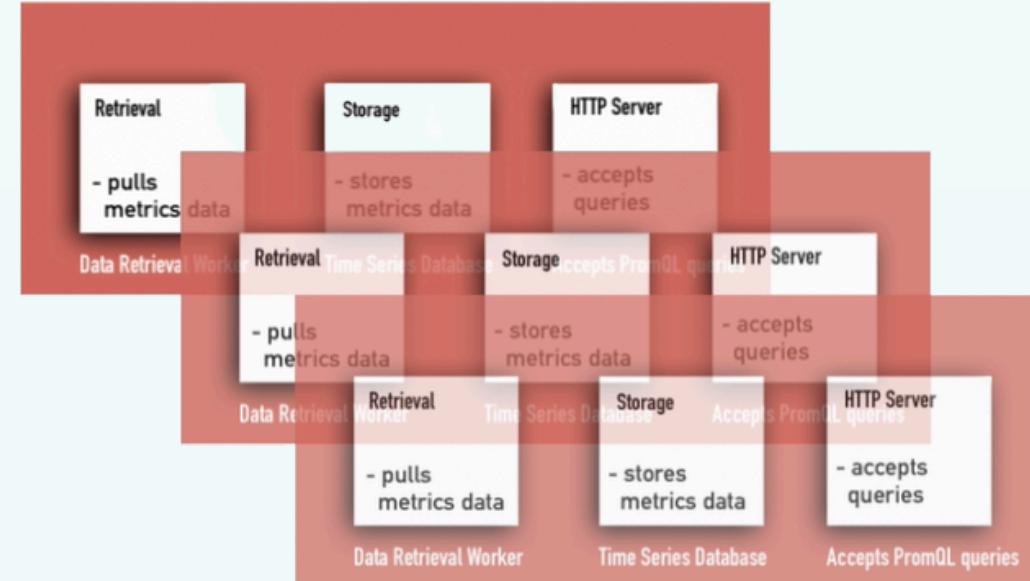
 Works, even if other parts of infrastructure broken

 No extensive set-up needed

 Less complex



Difficult to scale



Limits Monitoring

Workaround:



Increase Prometheus server capacity



Limit number of metrics

Other Prometheus Features

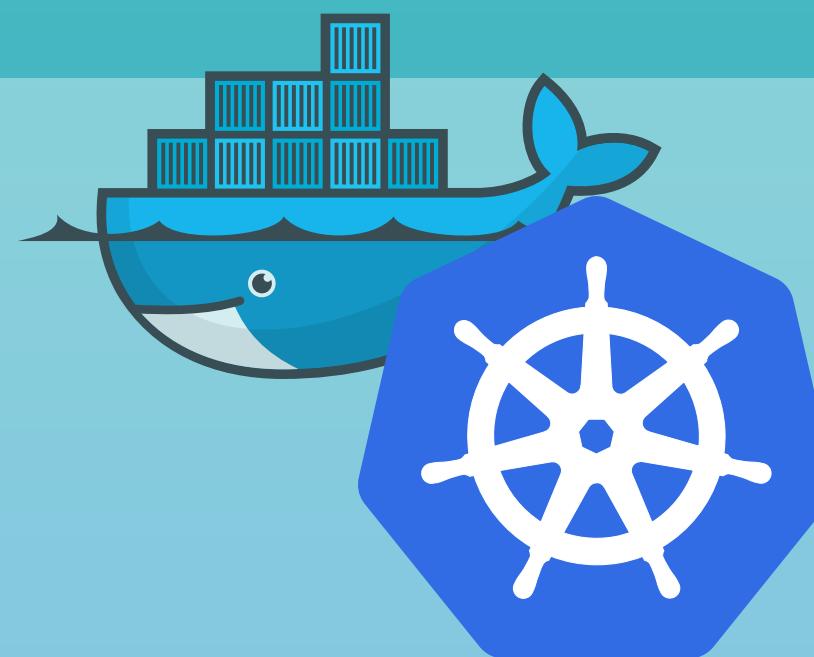
Prometheus Federation

- Allows Prometheus to **scale to environments with tens of data centers and millions of nodes**
- Allows a Prometheus server to scrape data from other Prometheus servers



Prometheus with Docker and Kubernetes

- Fully compatible
- Prometheus components available as Docker images
- Can easily be deployed in container environments like K8s
- **Monitoring of K8s cluster node resources out-of-the-box!**



Deploy Monitoring Stack - 1

3 different ways to deploy the Prometheus monitoring stack

1) Do it yourself

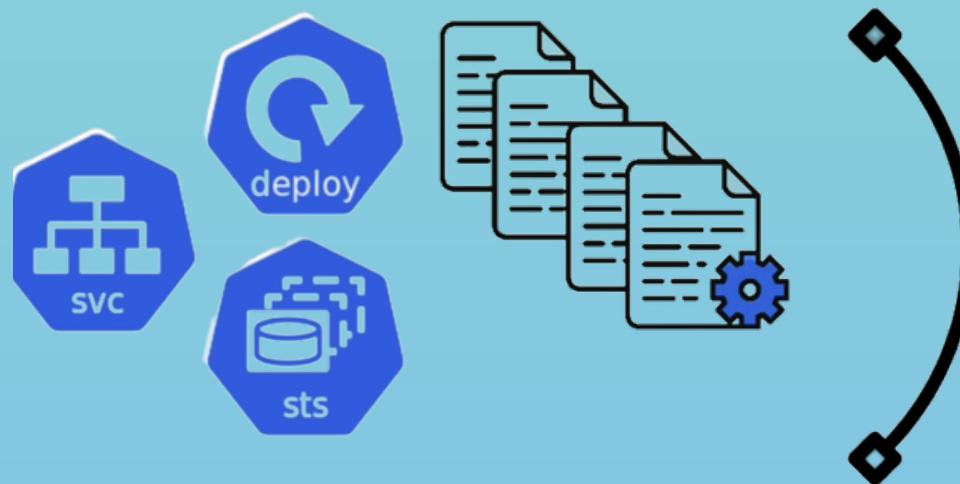
- Create all configuration **YAML files yourself**
- Execute them in right order



2) Using an Operator

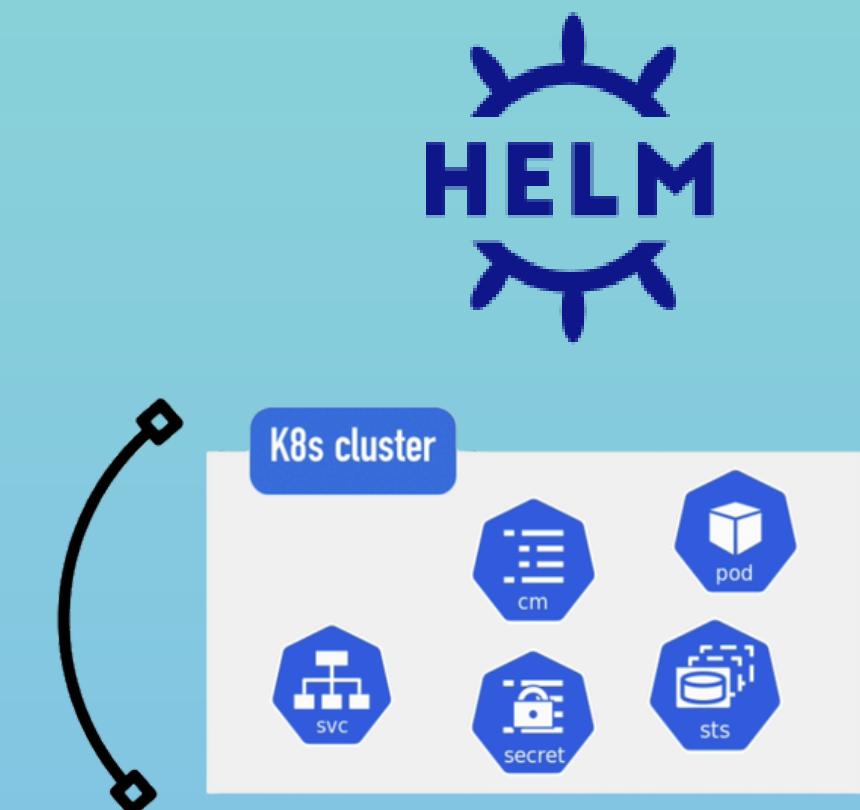
- Manager of all Prometheus components

1. Find Prometheus operator
2. Deploy in K8s cluster



3) Using Helm

- Using Helm chart to deploy operator
- Helm: Manage initial setup
- Operator: Manage setup



Deploy Monitoring Stack - 2

Overview of K8s resources deployed:

3 Deployments

- Prometheus Operator

created Prometheus and
Alertmanager StatefulSet

- Grafana
- Kube State Metrics

=> own Helm chart

=> dependency of this Helm chart

=> scrapes K8s components

1 DaemonSet

- Node Exporter DaemonSet

=> connects to server

=> translates Worker Nodes
metrics to Prometheus metrics

CPU usage load on server



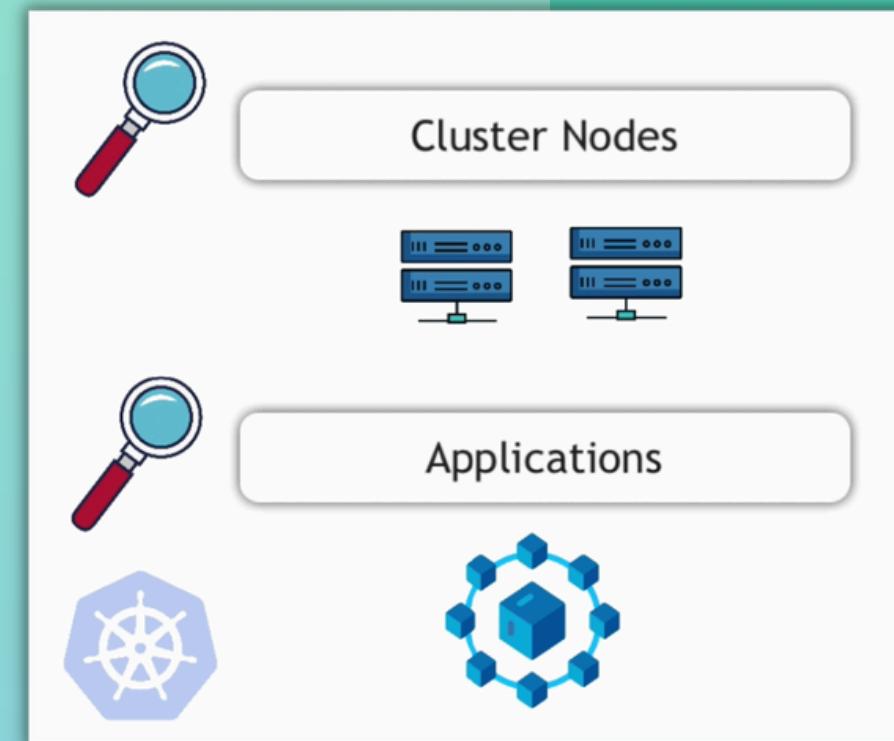
Data Visualization - 1

1st step: Decide what to monitor?

- Notice when something **unexpected** happens
- **Observe** any **anomalies**

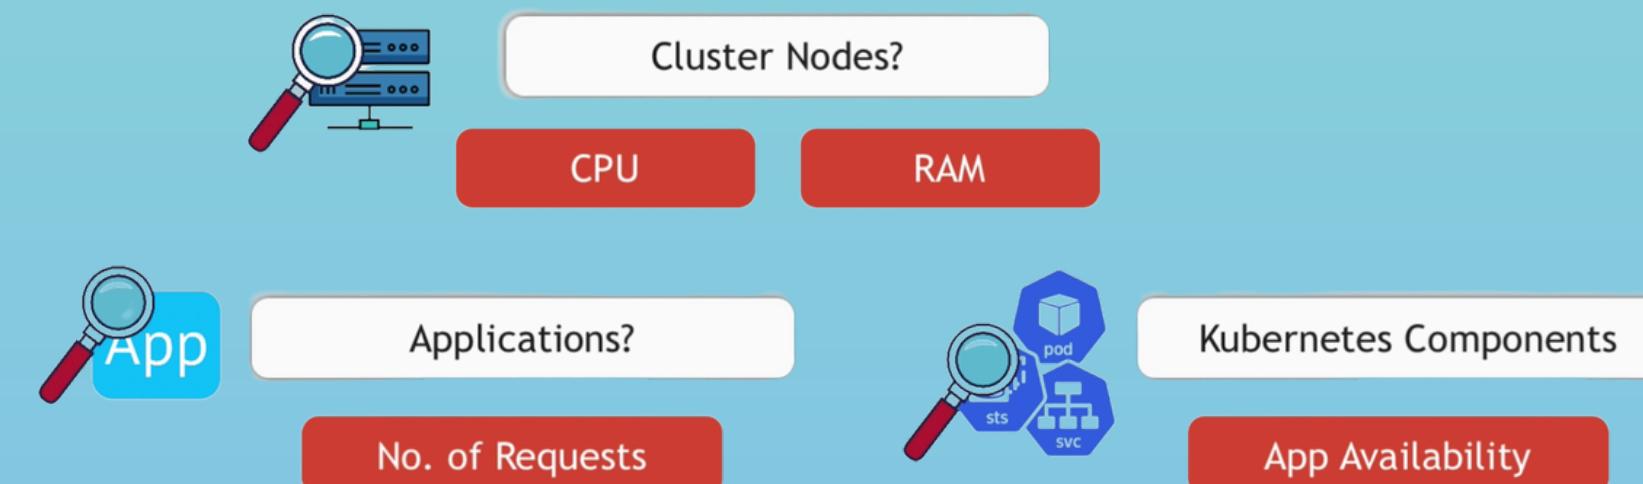
CPU spikes, insufficient storage, high load, unauthorized requests

- Analyze and **react** accordingly



2nd step: How to get this information?

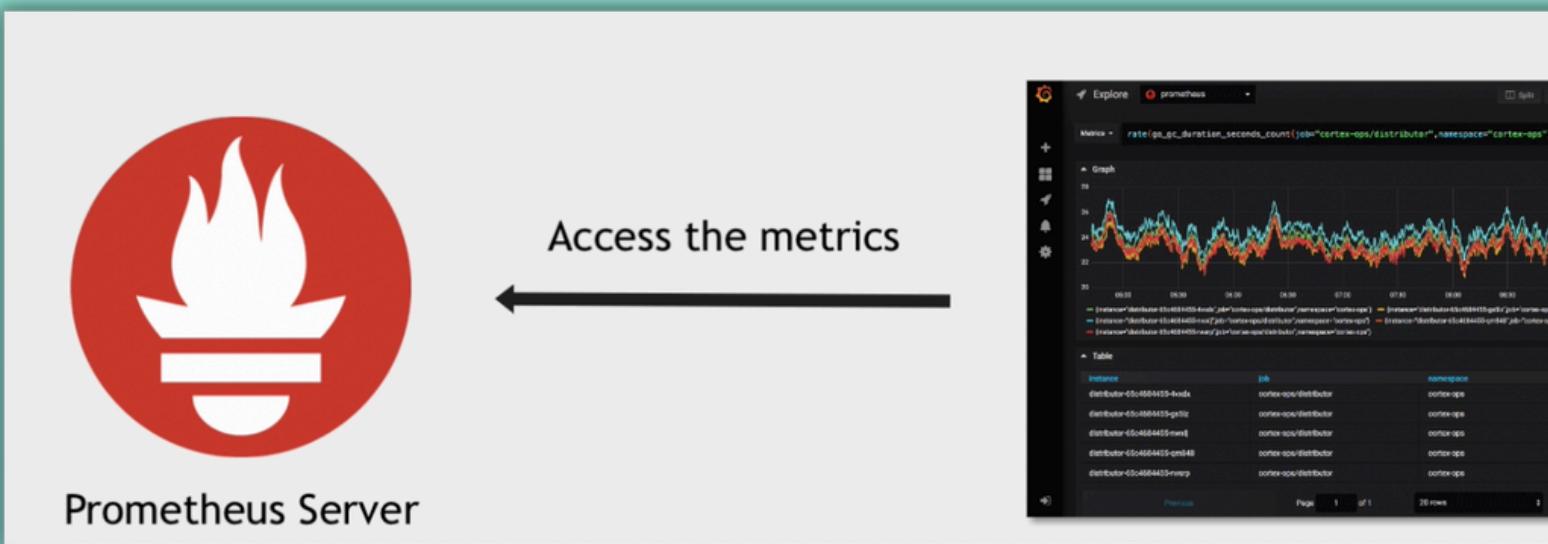
- How to get visibility of these monitoring data
- What data do we have available?



Data Visualization - 2

3rd step: Use a proper data visualization tool

- **Grafana** = a powerful open source visualization and analytics software
- Already deployed with the Prometheus Operator Helm Chart



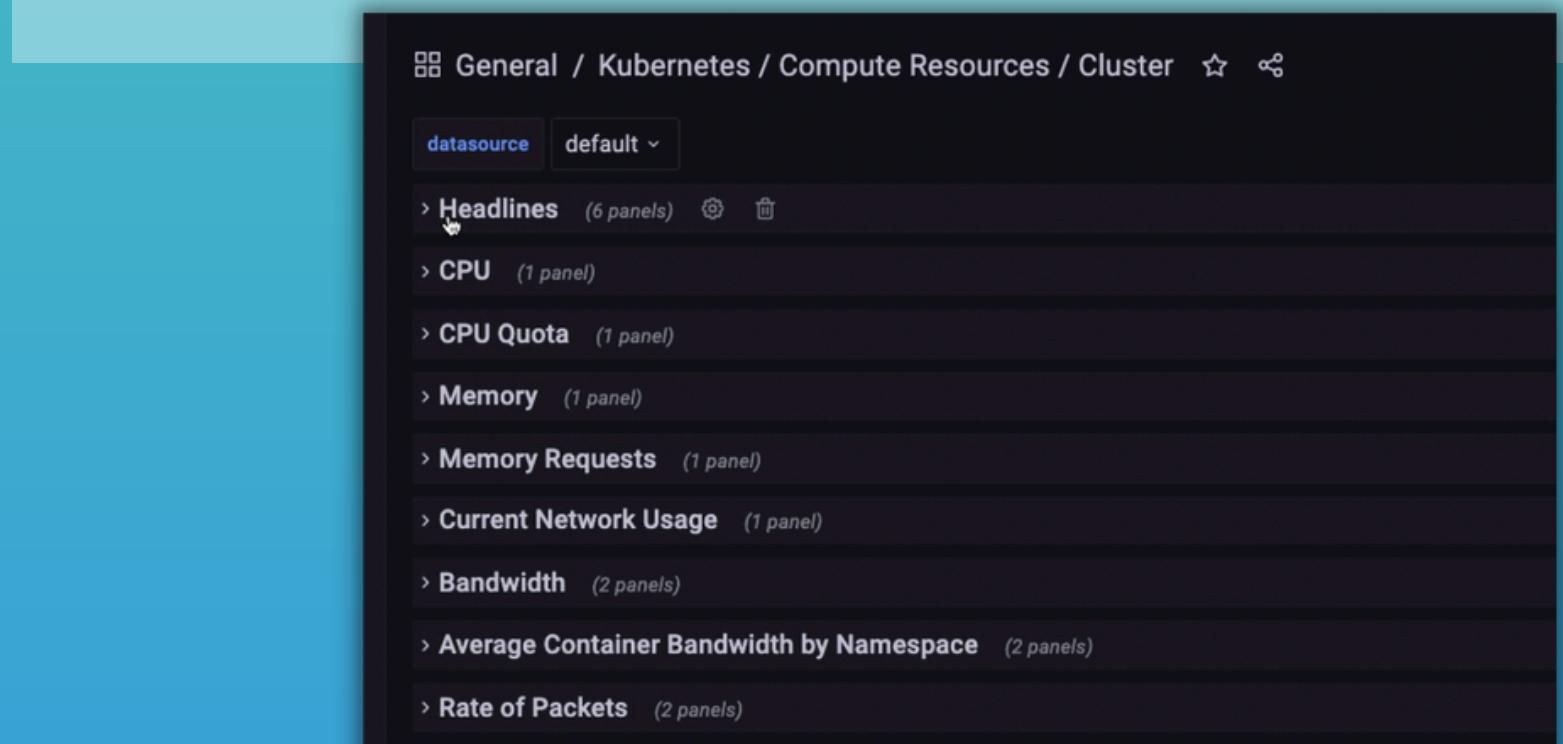
Data Visualization Tool

Grafana

With Grafana you can **create dynamic and reusable dashboards** that allow you to visualize your data in any way you want

Dashboard

- Dashboard is a set of one or more panels
- You can create your own Dashboards
- Organized into one or more rows
- Row is a logical divider within a dashboard
- Rows are used to group panels together



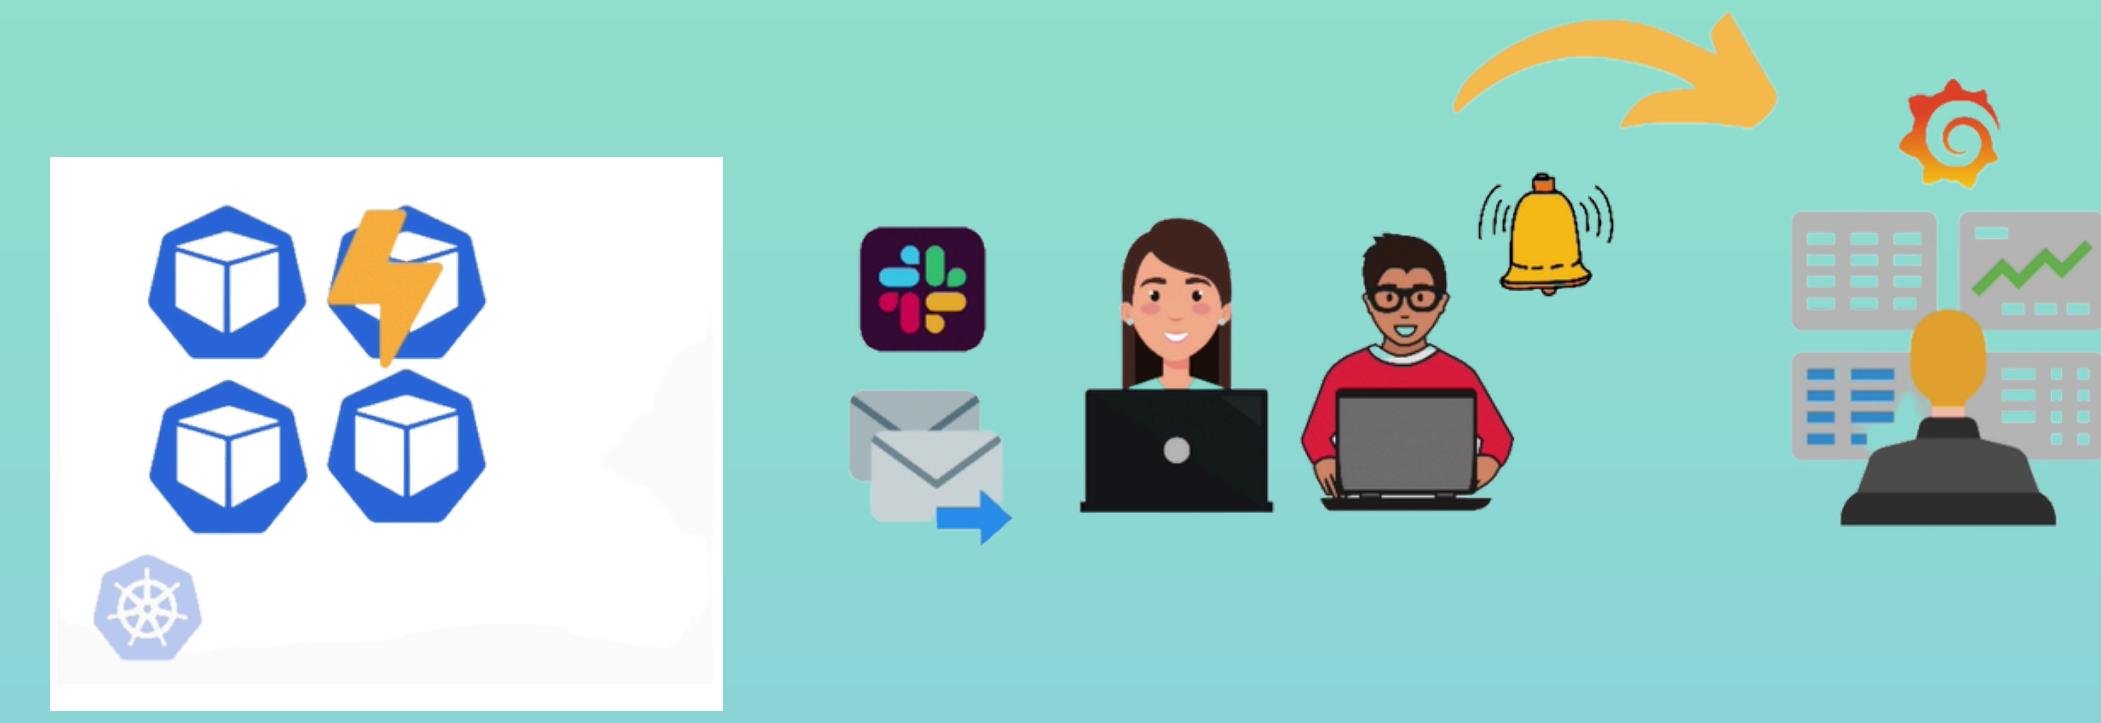
Panel

- The basic visualization building block in Grafana
- Composed by a query and a visualization
- Each panel has a query editor **specific to the data source** selected in the panel
- Can be moved and resized within a dashboard

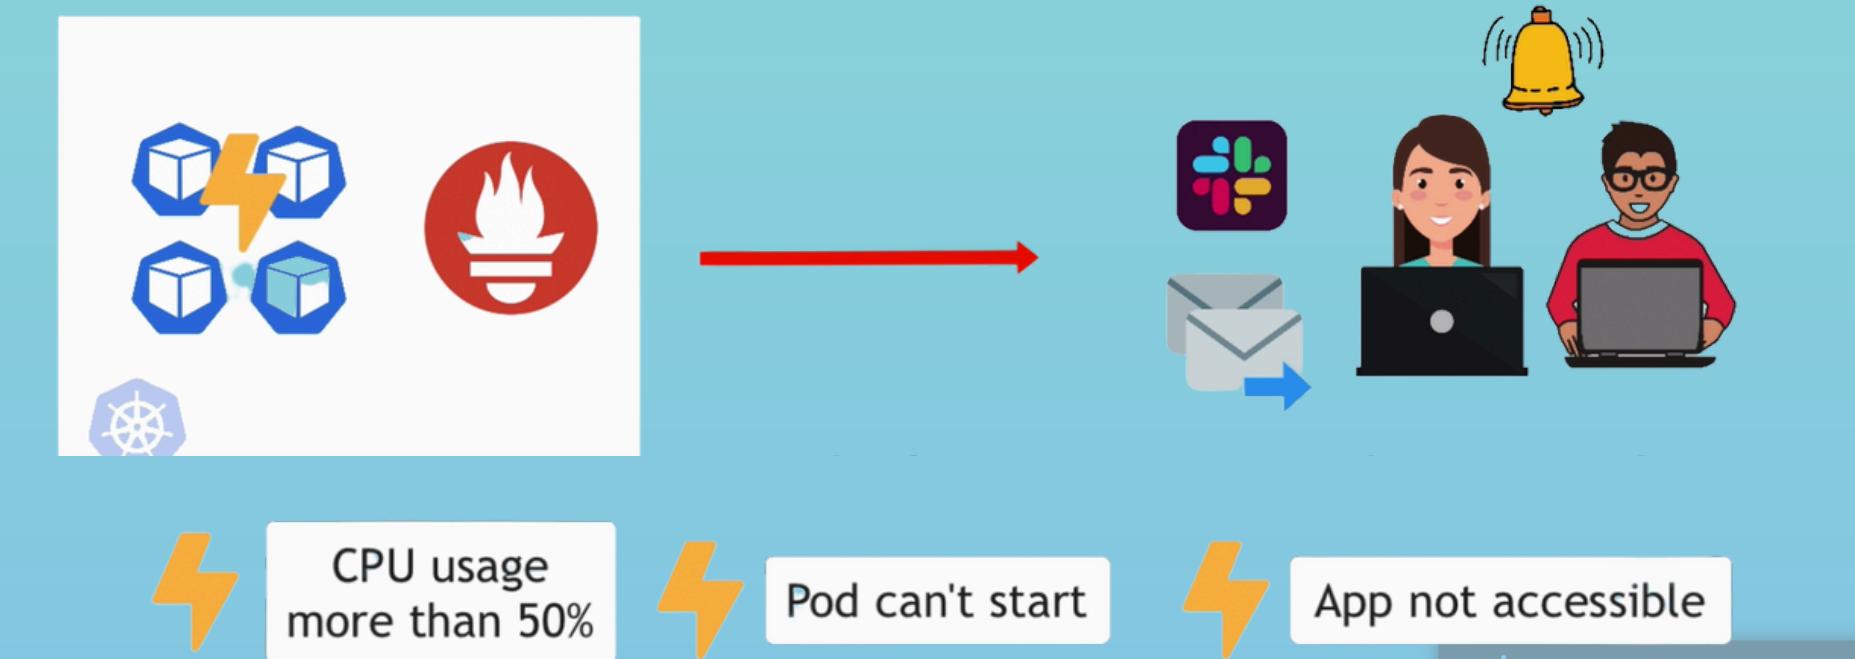


Alerting in Prometheus - 1

- Instead of constantly checking, you want to **get notified** when something happens
- Then you will check your dashboards



- For that we need to configure our monitoring stack to notify us whenever something unexpected happens



Alerting in Prometheus - 2

Configure Alerting

Alerting with Prometheus is separated **into 2 parts**:

- 1) **Alerting rules in Prometheus server** send alerts to an Alertmanager
- 2) **Alertmanager then manages (deduplicating, grouping, routing) those alerts**, including sending out notifications

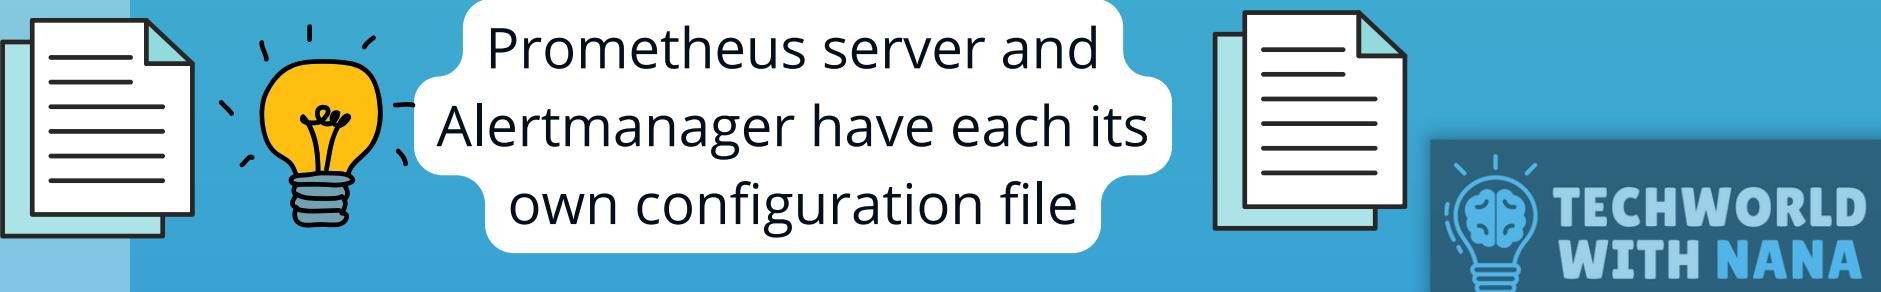
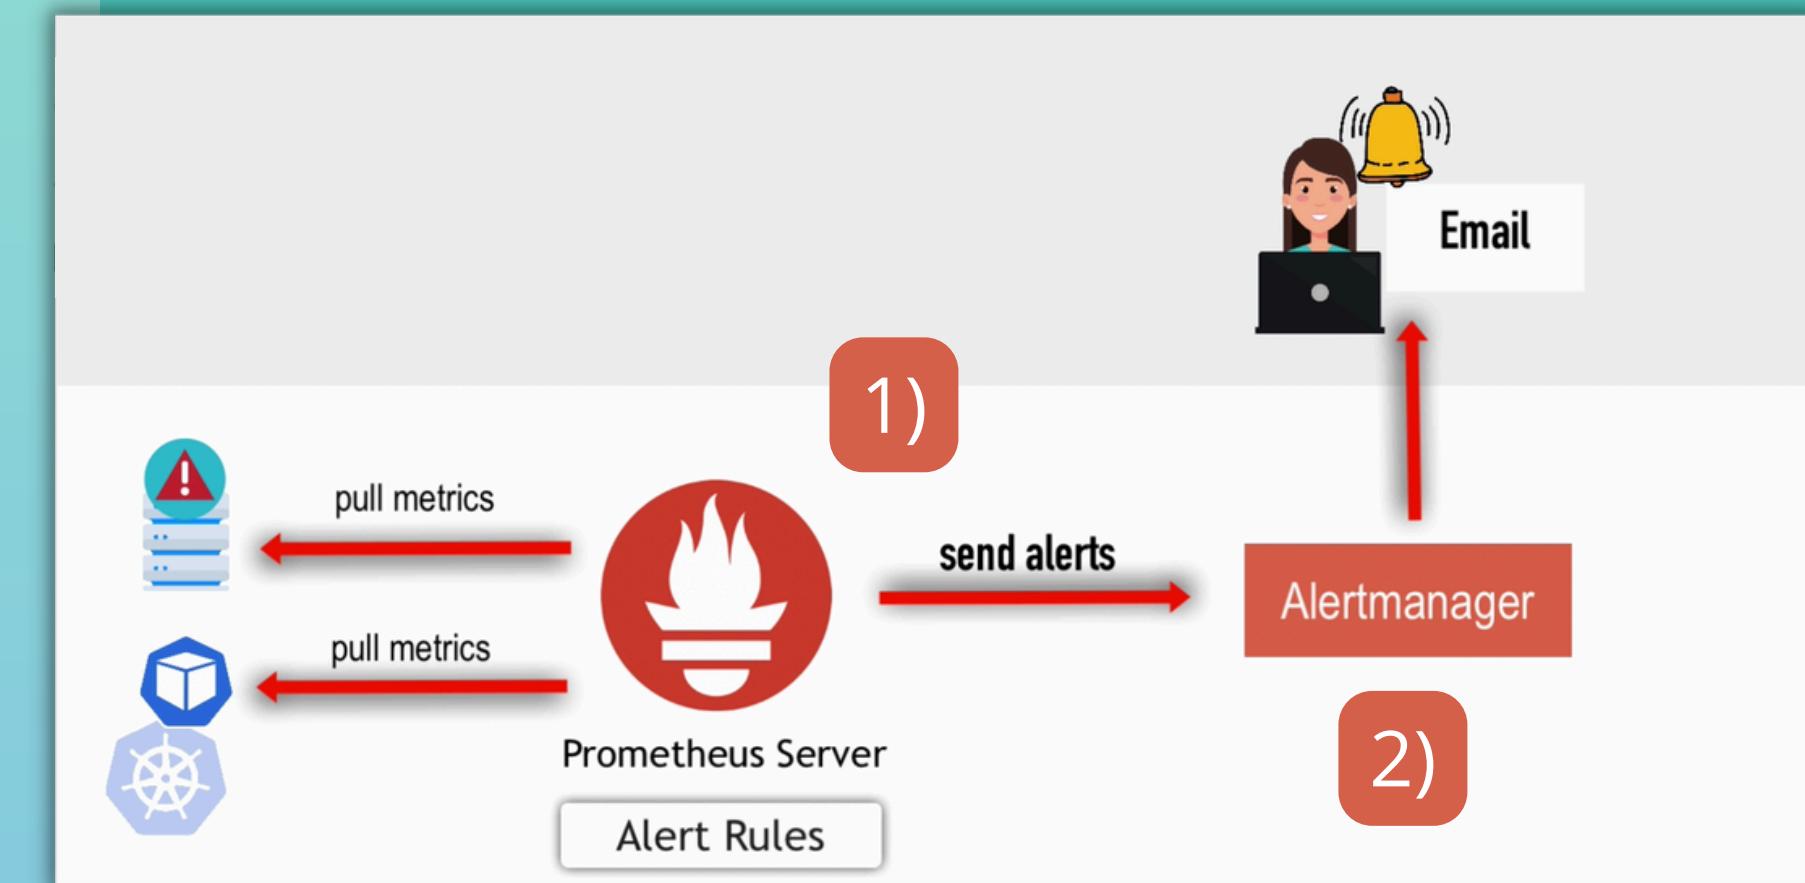
Main steps to setup alerting and notifications:

1. Setup and configure the Alertmanager
2. Configure Prometheus to talk to the Alertmanager
3. Create alerting rules in Prometheus

Example Alert rules to configure

1st Alert: when CPU usage > 50%

2nd Alert: when Pod cannot start



Alerting in Prometheus - 3

Alertmanager example configuration

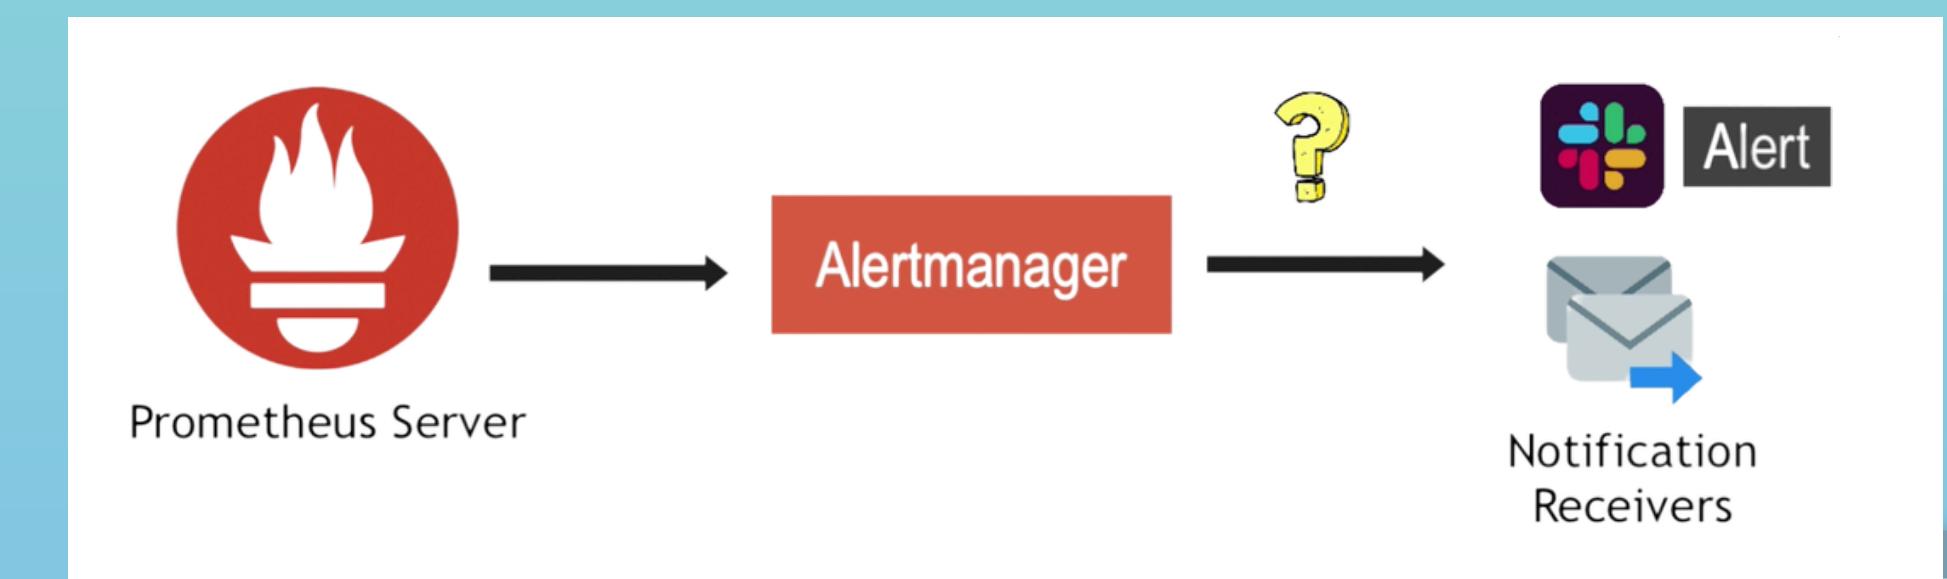
```
global:  
  resolve_timeout: 5m  
  http_config:  
    follow_redirects: true  
  smtp_hello: localhost  
  smtp_require_tls: true  
  pagerduty_url: https://events.pagerduty.com/v2/enqueue  
  opsgenie_api_url: https://api.opsgenie.com/  
  wechat_api_url: https://qyapi.weixin.qq.com/cgi-bin/  
  victorops_api_url: https://alert.victorops.com/integrations/g  
route:  
  receiver: "null"  
  group_by:  
  - job  
  continue: false  
  routes:  
  - receiver: "null"  
    match:  
      alertname: Watchdog  
    continue: false  
  group_wait: 30s  
  group_interval: 5m  
  repeat_interval: 12h  
receivers:  
- name: "null"  
templates:
```

Any Alert

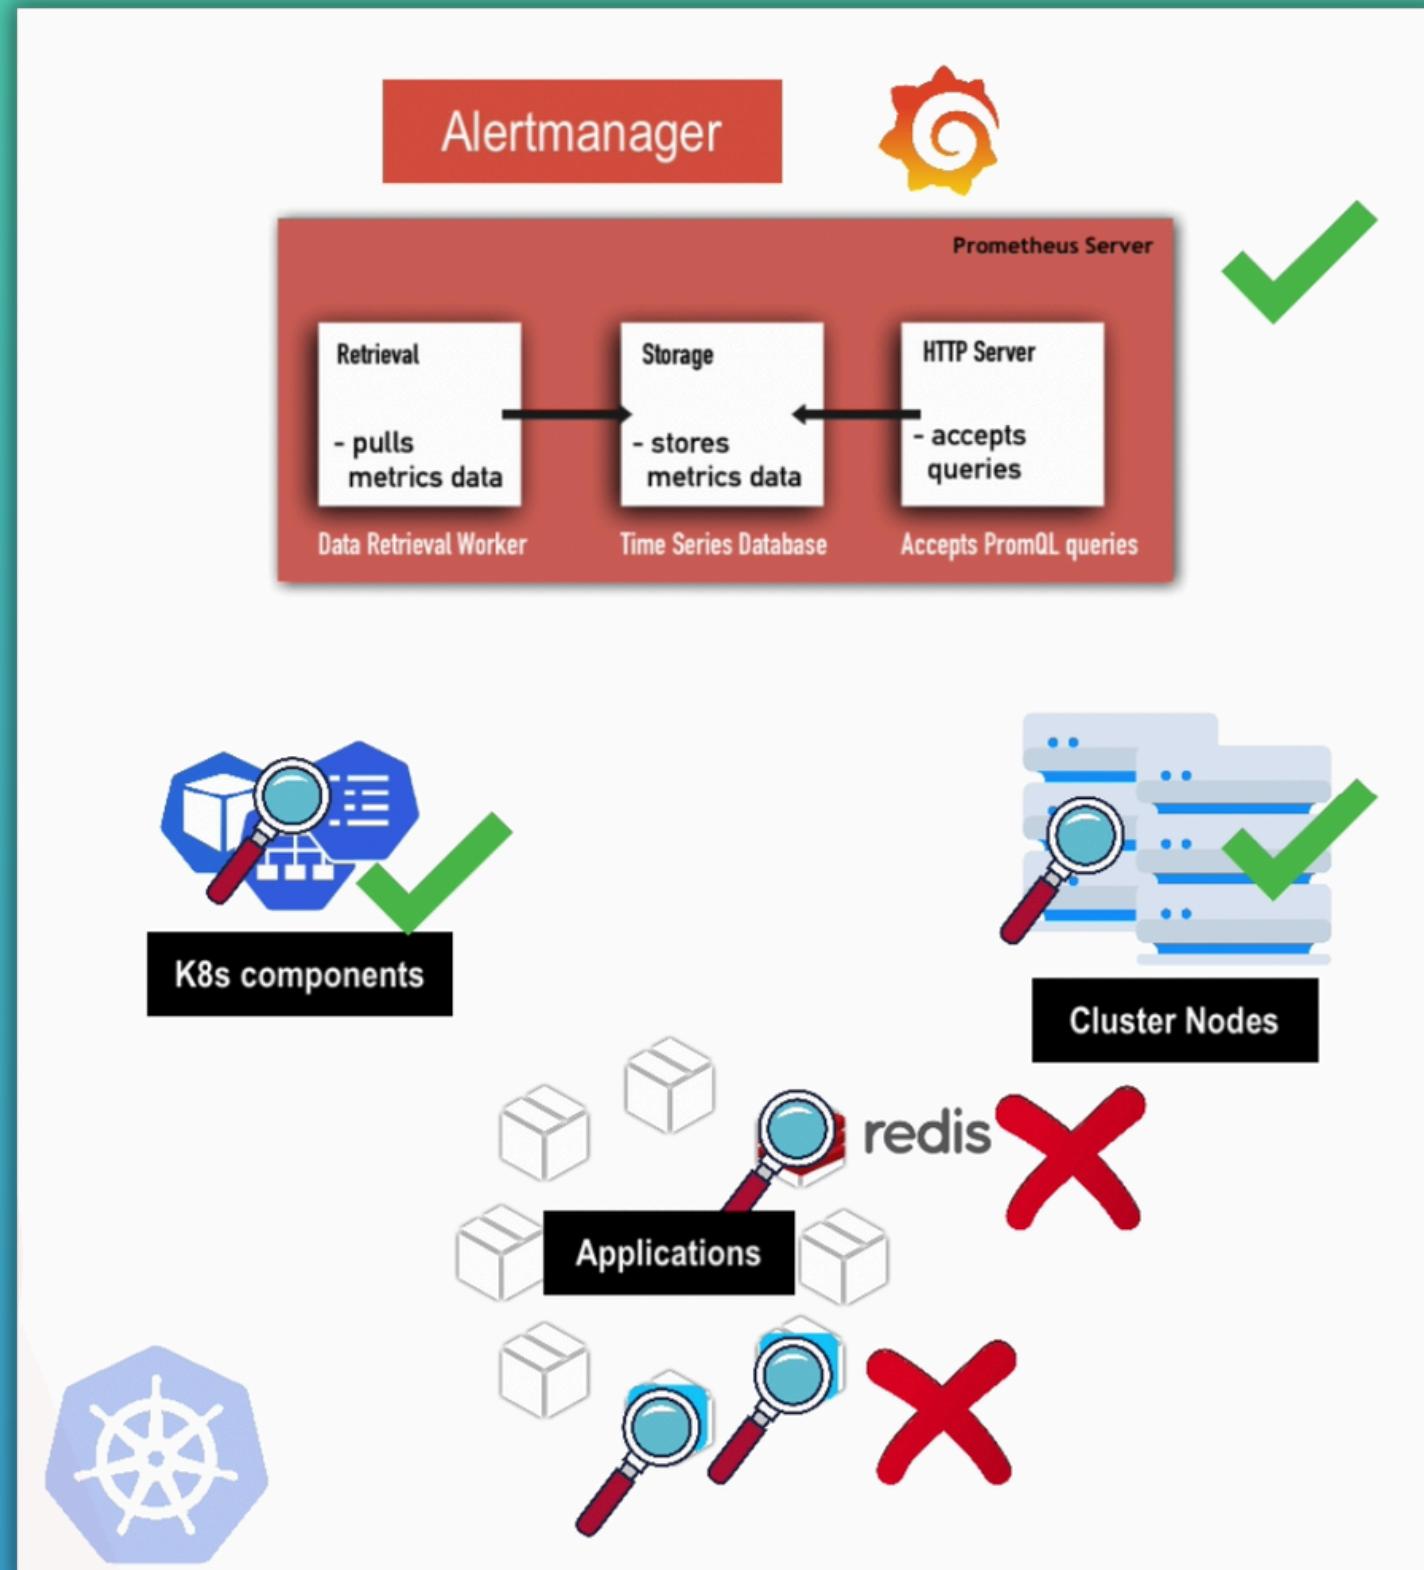
Specific Alerts

Receiver:

- These are the notification integrations
- For each alert you can define own receiver. For example:
 - send all K8s cluster related issues to admin email
 - send all application related issues to developer team's slack channel



Monitor third party and own applications - 1



Still missing:

Configure Third-Party and own application monitoring

- ✓ Monitor Kubernetes components
 - ✓ Monitor Resource Consumption on the Nodes
 - ✓ Monitor Prometheus Stack itself
 - ✗ Monitor third-party applications like Redis
 - ✗ Monitor own applications, like your online shop
- microservices

Monitor third party and own applications - 2

3rd-party example: Redis

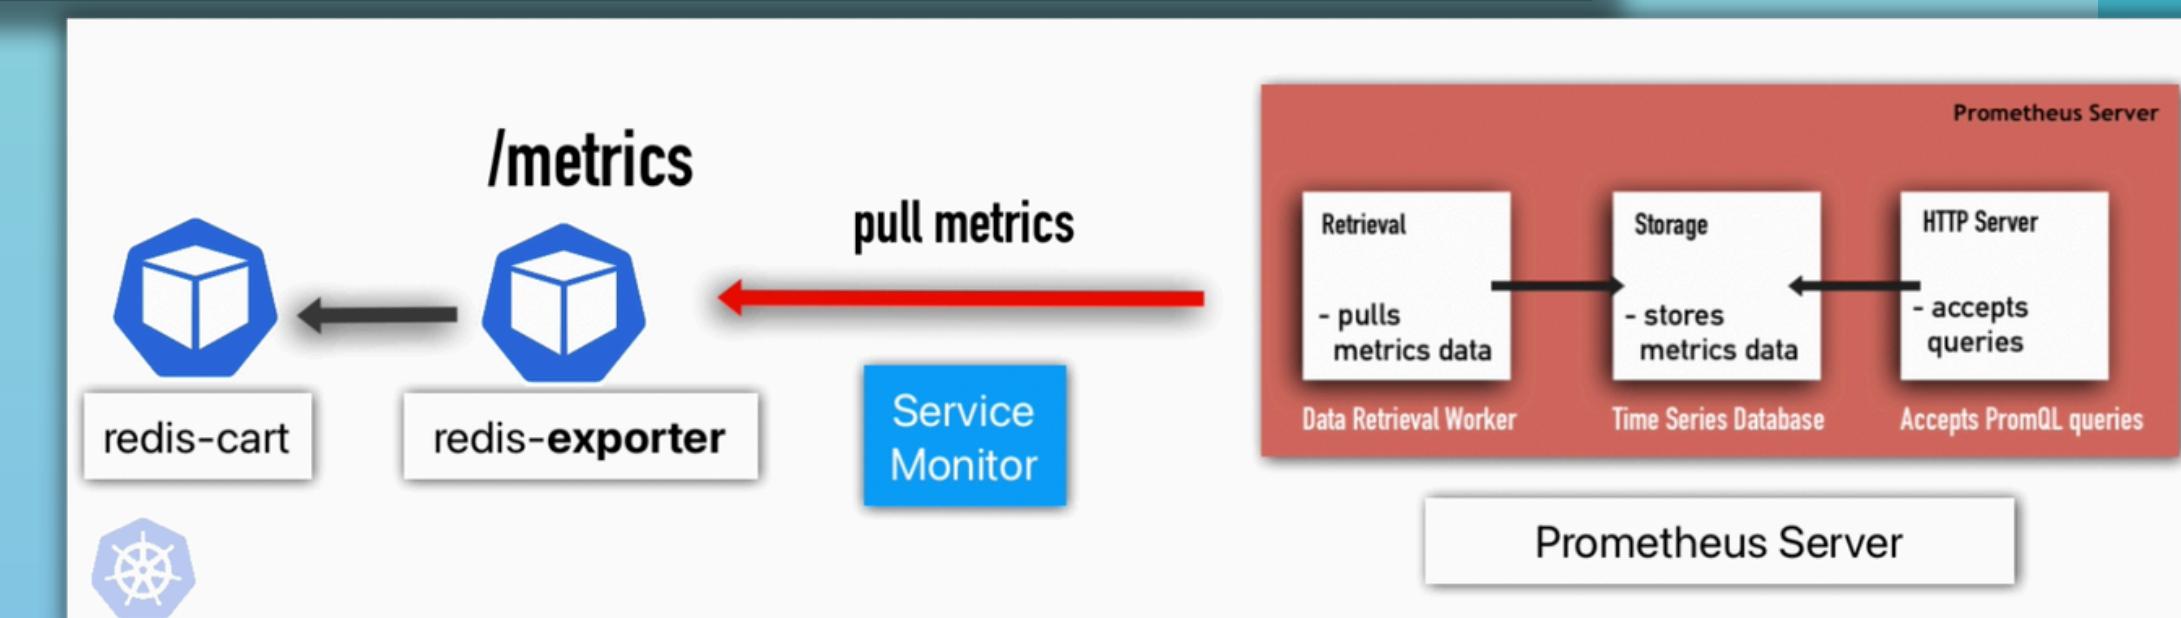
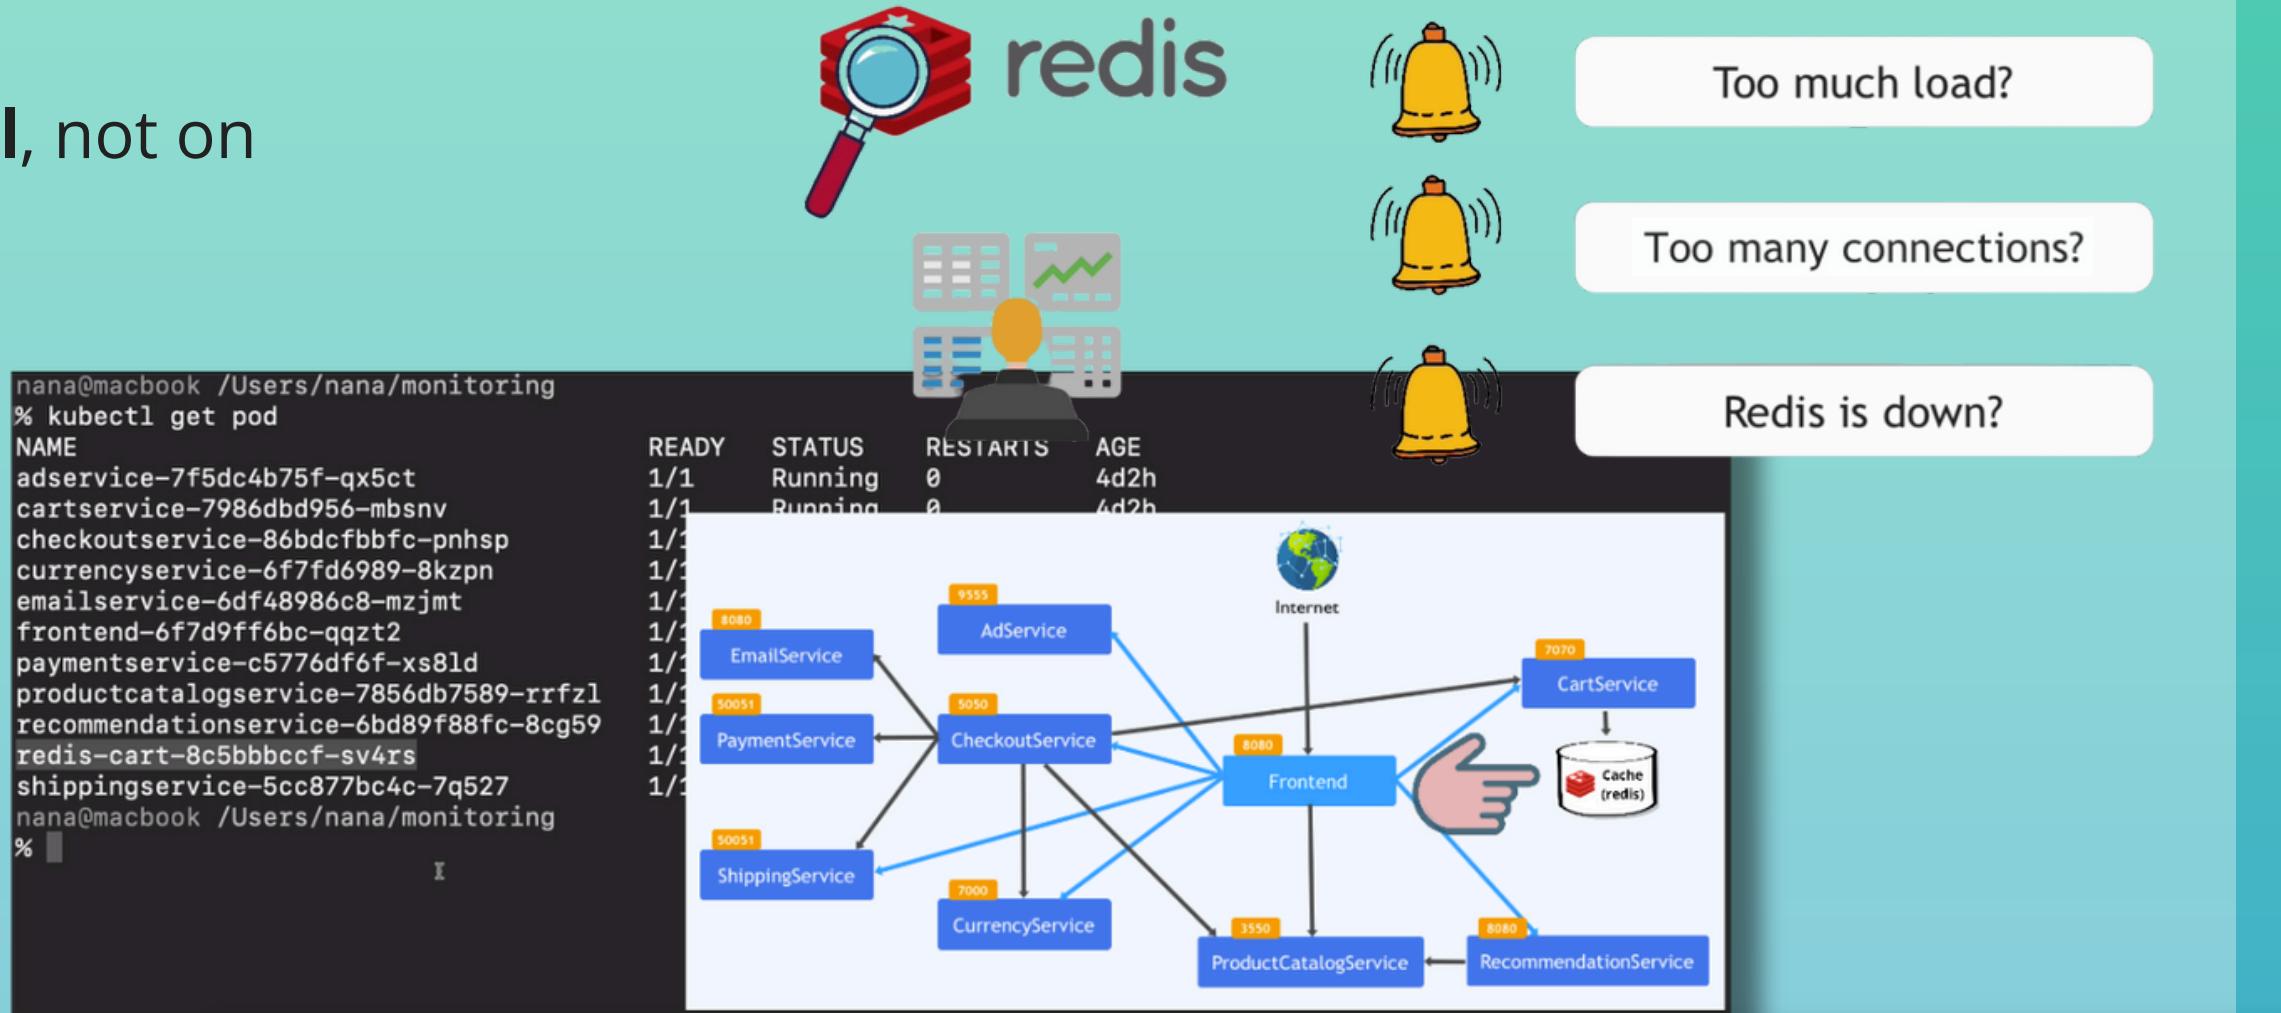
- Monitor Redis on application level, not on Kubernetes level



As we learnt, we can do that via an **Exporter**!

How to:

1. Deploy redis-exporter
2. Deploy ServiceMonitor (custom K8s resource) to tell Prometheus about this new exporter



Monitor third party and own applications - 3

Own application

- No exporter available for your own application
- So we have to define the metrics ourselves

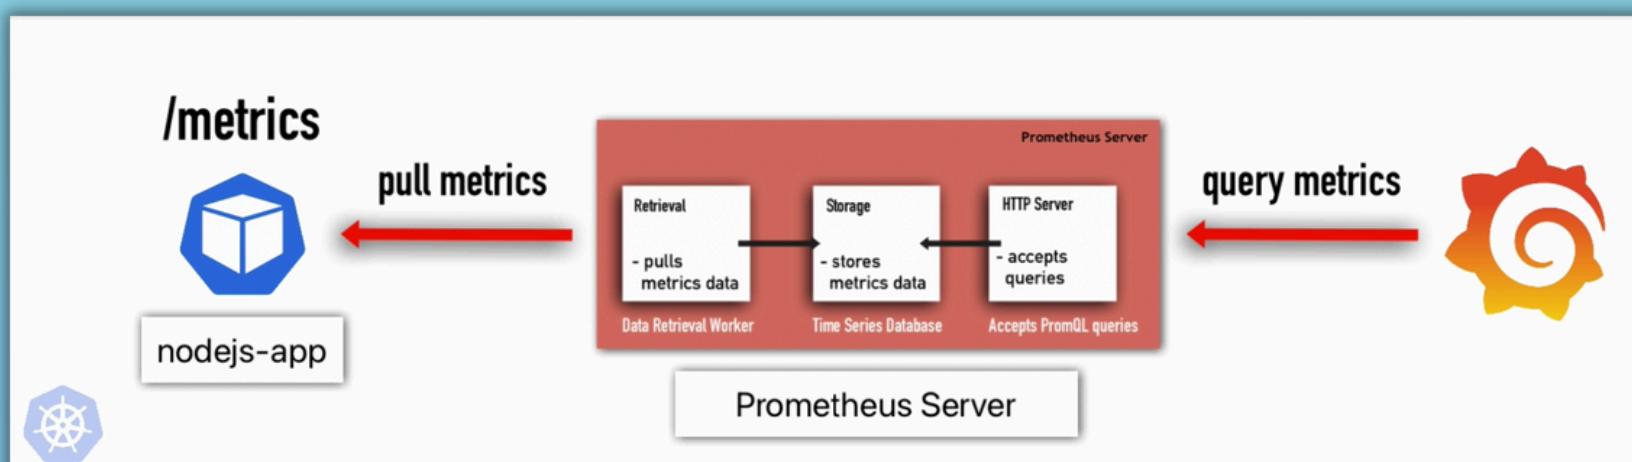


How to (Nodejs application):

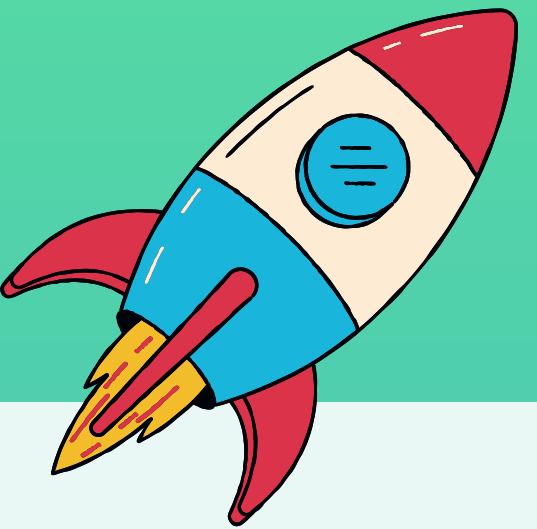
1. Expose metrics using Nodejs client library
2. Deploy Nodejs application in the cluster
3. Configure Prometheus to scrape new target
(ServiceMonitor)
4. Visualize scraped metrics in Grafana Dashboard

Client Libraries:

- Gives you an abstract interface to expose your metrics
- Libraries implement the Prometheus metric types: Counter, Gauge, Histogram, Summary
- Choose client library that matches the application's language



Best Practices



Official Best Practices:

- Metric and Label Naming: <https://prometheus.io/docs/practices/naming/>
- Set of guidelines for instrumenting your code:
<https://prometheus.io/docs/practices/instrumentation/>
- Consoles and Dashboards: <https://prometheus.io/docs/practices/consoles/>
- Alerting: <https://prometheus.io/docs/practices/alerting/>
- On when to use the Pushgateway: <https://prometheus.io/docs/practices/pushing/>